

Train a Neural Networks using Keras

Shusen Wang

Implement a Softmax Classifier Using Keras

Deep Learning Systems



Softmax Classifier

The MNIST Dataset

5	0	4	1	9	2	1	3	1	4
3	5	3	6	1	7	2	8	6	9
4	0	9	1	1	2	4	3	2	7
3	8	6	9	0	5	6	0	7	6
1	8	7	9	3	9	8	5	3	3
3	0	7	4	9	8	0	9	4	1
4	4	6	0	4	5	6	1	0	0
1	7	1	6	3	0	2	1	1	7
9	0	2	6	7	8	3	9	0	4
6	7	4	6	8	0	7	8	3	1

- $n = 60,000$ training samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$.
- Each \mathbf{x}_j is a 28×28 image.
- Each y_j is an integer in $\{0, 1, 2, \dots, 9\}$.

Softmax Classifier

The MNIST Dataset

```
5 0 4 1 9 2 1 3 1 4  
3 5 3 6 1 7 2 8 6 9  
4 0 9 1 1 2 4 3 2 7  
3 8 6 9 0 5 6 0 7 6  
1 8 7 9 3 9 8 5 3 3  
3 0 7 4 9 8 0 9 4 1  
4 4 6 0 4 5 6 1 0 0  
1 7 1 6 3 0 2 1 1 7  
9 0 2 6 7 8 3 9 0 4  
6 7 4 6 8 0 7 8 3 1
```

- $n = 60,000$ training samples $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$.
- Each \mathbf{x}_j is a 28×28 image.
- Each y_j is an integer in $\{0, 1, 2, \dots, 9\}$.

Softmax classifier ($\mathbf{f}: \mathbb{R}^{784} \mapsto \mathbb{R}^{10}$) for MNIST:

- **Input:** vector $\mathbf{x} \in \mathbb{R}^{784}$.
- $\mathbf{z} = \mathbf{W} \mathbf{x} + \mathbf{b} \in \mathbb{R}^{10}$.
- **Output:** $\mathbf{f}(\mathbf{x}) = \text{SoftMax}(\mathbf{z})$.

Trainable parameters:

- $\mathbf{W} \in \mathbb{R}^{10 \times 784}$
- $\mathbf{b} \in \mathbb{R}^{10}$

1. Load and Process the MNIST Dataset

Load data

```
from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

print('Shape of x_train: ' + str(x_train.shape))
print('Shape of x_test: ' + str(x_test.shape))
print('Shape of y_train: ' + str(y_train.shape))
print('Shape of y_test: ' + str(y_test.shape))
```

```
Shape of x_train: (60000, 28, 28)
Shape of x_test: (10000, 28, 28)
Shape of y_train: (60000,)
Shape of y_test: (10000,)
```

1. Load and Process the MNIST Dataset

Vectorization: convert the 28×28 images to 784-dim vectors.

```
x_train_vec = x_train.reshape(60000, 784)
x_test_vec = x_test.reshape(10000, 784)

print('Shape of x_train_vec is ' + str(x_train_vec.shape))

Shape of x_train_vec is (60000, 784)
```

1. Load and Process the MNIST Dataset

Vectorization: convert the **labels** (an integer in $\{0, 1, \dots, 9\}$) to 10-dim vectors (by the one-hot encode).

```
def to_one_hot(labels, dimension=10):
    results = np.zeros((len(labels), dimension))
    for i, label in enumerate(labels):
        results[i, label] = 1.
    return results

y_train_vec = to_one_hot(y_train)
y_test_vec = to_one_hot(y_test)
print('Shape of y_train_vec is ' + str(y_train_vec.shape))

Shape of y_train_vec is (60000, 10)
```

1. Load and Process the MNIST Dataset

Partition the **training** set to **training** and **validation** sets.

```
rand_indices = np.random.permutation(60000)
train_indices = rand_indices[0:50000]
valid_indices = rand_indices[50000:60000]

x_valid_vec = x_train_vec[valid_indices, :]
y_valid_vec = y_train_vec[valid_indices, :]

x_train_vec = x_train_vec[train_indices, :]
y_train_vec = y_train_vec[train_indices, :]

print('Shape of x_valid_vec: ' + str(x_valid_vec.shape))
print('Shape of y_valid_vec: ' + str(y_valid_vec.shape))
print('Shape of x_train_vec: ' + str(x_train_vec.shape))
print('Shape of y_train_vec: ' + str(y_train_vec.shape))

Shape of x_valid_vec: (10000, 784)
Shape of y_valid_vec: (10000, 10)
Shape of x_train_vec: (50000, 784)
Shape of y_train_vec: (50000, 10)
```

2. Build the Softmax Classifier

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(10, activation='softmax', input_shape=(784,)))

# print the summary of the model.
model.summary()
```

$$f(x) = \text{SoftMax}(Wx + b)$$



Layer (type)	Output Shape	Param #
<hr/>		
dense_1 (Dense)	(None, 10)	7850
<hr/>		
Total params:	7,850	
Trainable params:	7,850	
Non-trainable params:	0	

3. Train the Softmax Classifier

Specify: optimization algorithm, learning rate (LR), loss function, and metric.

```
from keras import optimizers  
model.compile(optimizers.RMSprop(lr=0.0001),  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

3. Train the Softmax Classifier

Specify: batch size and number of epochs.

```
history = model.fit(x_train_vec, y_train_vec,
                     batch_size=128, epochs=50,
                     validation_data=(x_valid_vec, y_valid_vec))
```

```
Train on 50000 samples, validate on 10000 samples
Epoch 1/50
50000/50000 [=====] - 1s 19us/step - loss: 11.2824 - acc: 0.2816 - val_loss: 8.7464 - val_acc: 0.4398
Epoch 2/50
50000/50000 [=====] - 1s 13us/step - loss: 7.1525 - acc: 0.5384 - val_loss: 5.7830 - val_acc: 0.6221
Epoch 3/50
50000/50000 [=====] - 1s 13us/step - loss: 5.1013 - acc: 0.6661 - val_loss: 4.5210 - val_acc: 0.7036
Epoch 4/50
50000/50000 [=====] - 1s 13us/step - loss: 4.0385 - acc: 0.7339 - val_loss: 3.7206 - val_acc: 0.7533
...
●
●
●
Epoch 49/50
50000/50000 [=====] - 1s 13us/step - loss: 1.1430 - acc: 0.9228 - val_loss: 1.3567 - val_acc: 0.9088
Epoch 50/50
50000/50000 [=====] - 1s 14us/step - loss: 1.1407 - acc: 0.9230 - val_loss: 1.3526 - val_acc: 0.9077
```

4. Examine the Results

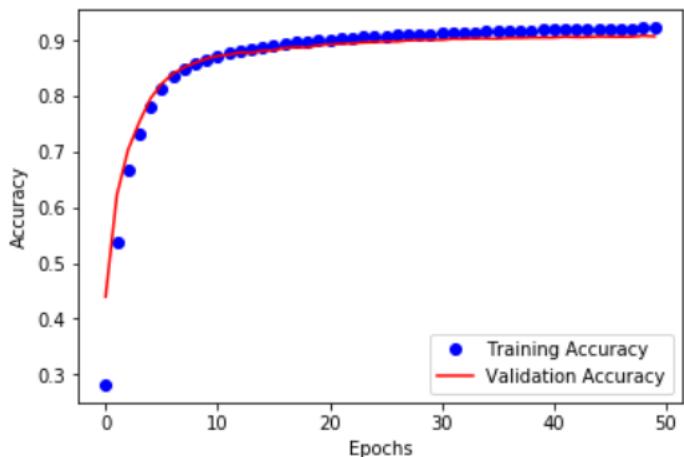
Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).

```
import matplotlib.pyplot as plt
%matplotlib inline

epochs = range(50) # 50 is the number of epochs
train_acc = history.history['acc']
valid_acc = history.history['val_acc']
plt.plot(epochs, train_acc, 'bo', label='Training Accuracy')
plt.plot(epochs, valid_acc, 'r', label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

4. Examine the Results

Plot the **accuracy** against **epochs** (1 epoch = 1 pass over the data).



Why using the validation set?

- Tune the hyper-parameters, e.g., learning rate, regularization.
- Check whether overfit or underfit.

4. Examine the Results

Evaluate the model on the ***test*** set. (So far, the model has not seen the ***test*** set.)

```
loss_and_acc = model.evaluate(x_test_vec, y_test_vec)
print('loss = ' + str(loss_and_acc[0]))
print('accuracy = ' + str(loss_and_acc[1]))

10000/10000 [=====] - 0s 20us/step
loss = 1.2781493856459185
accuracy = 0.9131
```

4. Examine the Results

- Training accuracy: 92.3%
- Validation accuracy: 90.8%
- Test accuracy: 91.3%

Not bad! A random guess has only 10% accuracy.

Implement a Neural Network Using Keras

Full-Connected Neural Network

- **Input:** vector $\mathbf{x}^{(0)} \in \mathbb{R}^{784}$.

- $\mathbf{z}^{(1)} = \mathbf{W}^{(0)} \mathbf{x}^{(0)} + \mathbf{b}^{(0)} \in \mathbb{R}^{d_1}$.
- $\mathbf{x}^{(1)} = \max\{\mathbf{0}, \mathbf{z}^{(1)}\} \in \mathbb{R}^{d_1}$.

Hidden Layer 1

- $\mathbf{z}^{(2)} = \mathbf{W}^{(1)} \mathbf{x}^{(1)} + \mathbf{b}^{(1)} \in \mathbb{R}^{d_2}$.
- $\mathbf{x}^{(2)} = \max\{\mathbf{0}, \mathbf{z}^{(2)}\} \in \mathbb{R}^{d_2}$.

Hidden Layer 2

- $\mathbf{z}^{(3)} = \mathbf{W}^{(2)} \mathbf{x}^{(2)} + \mathbf{b}^{(2)} \in \mathbb{R}^{10}$.
- **Output:** SoftMax($\mathbf{z}^{(3)}$).

Output Layer

Trainable parameters:

- $\mathbf{W}^{(0)} \in \mathbb{R}^{d_1 \times 784}$,
- $\mathbf{W}^{(1)} \in \mathbb{R}^{d_2 \times d_1}$,
- $\mathbf{W}^{(2)} \in \mathbb{R}^{10 \times d_2}$,
- $\mathbf{b}^{(0)} \in \mathbb{R}^{d_1}$,
- $\mathbf{b}^{(1)} \in \mathbb{R}^{d_1}$,
- $\mathbf{b}^{(2)} \in \mathbb{R}^{10}$.

2. Build the Softmax Classifier

```
from keras import models
from keras import layers

d1 = 500 # width of the 1st hidden layer
d2 = 500 # width of the 2nd hidden layer

model = models.Sequential()
model.add(layers.Dense(d1, activation='relu', input_shape=(784,)))
model.add(layers.Dense(d2, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

# print the summary of the model.
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 500)	392500
dense_5 (Dense)	(None, 500)	250500
dense_6 (Dense)	(None, 10)	5010
=====		
Total params:	648,010	
Trainable params:	648,010	
Non-trainable params:	0	

3. Train the Softmax Classifier

Specify: optimization algorithm, learning rate (LR), loss function, and metric.

```
from keras import optimizers  
model.compile(optimizers.RMSprop(lr=0.0001),  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

3. Train the Softmax Classifier

Specify: batch size and number of epochs.

```
history = model.fit(x_train_vec, y_train_vec,  
                     batch_size=128, epochs=50,  
                     validation_data=(x_valid_vec, y_valid_vec))
```

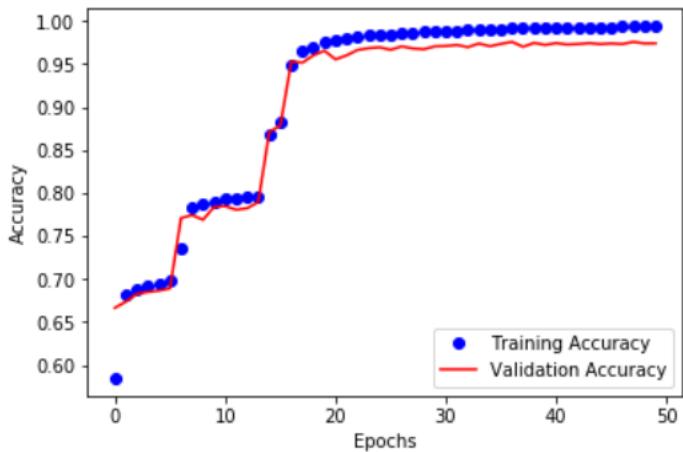
```
Train on 50000 samples, validate on 10000 samples  
Epoch 1/50  
50000/50000 [=====] - 5s 92us/step - loss: 6.5323 - acc: 0.5850 - val_loss: 5.2988 - val_acc: 0.6666  
Epoch 2/50  
50000/50000 [=====] - 4s 84us/step - loss: 5.0559 - acc: 0.6817 - val_loss: 5.1904 - val_acc: 0.6738  
Epoch 3/50  
50000/50000 [=====] - 3s 64us/step - loss: 4.9762 - acc: 0.6881 - val_loss: 5.0556 - val_acc: 0.6828
```

●
●
●

```
Epoch 49/50  
50000/50000 [=====] - 5s 93us/step - loss: 0.0912 - acc: 0.9934 - val_loss: 0.3404 - val_acc: 0.9738  
Epoch 50/50  
50000/50000 [=====] - 4s 79us/step - loss: 0.0920 - acc: 0.9934 - val_loss: 0.3327 - val_acc: 0.9739
```

4. Examine the Results

Plot the *accuracy* against *epochs* (1 epoch = 1 pass over the data).



4. Examine the Results

Evaluate the model on the ***test*** set. (So far, the model has not seen the ***test*** set.)

```
loss_and_acc = model.evaluate(x_test_vec, y_test_vec)
print('loss = ' + str(loss_and_acc[0]))
print('accuracy = ' + str(loss_and_acc[1]))
```

10000/10000 [=====] - 0s 43us/step
loss = 0.3648844491034643
accuracy = 0.9724

4. Examine the Results

- Training accuracy: 99.3%
- Validation accuracy: 97.4%
- Test accuracy: 97.2%

Much better than the simple softmax classifier (91.3% accuracy).

Summary