

# Variational Autoencoder for Image Generation

Shusen Wang

# Preliminary: Distance between Two Probability Distributions

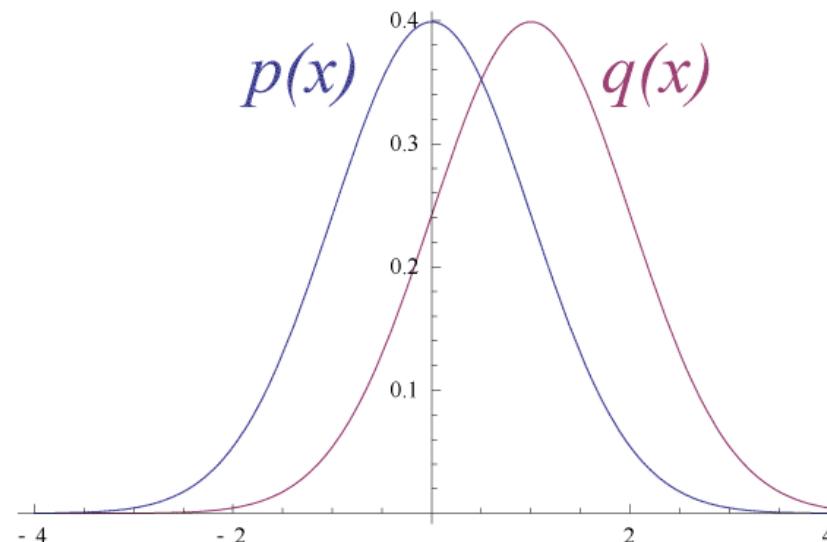
# Distance between Two Vectors

- Distance between two vectors is measured by vector norms.
- Examples:
  - $\ell_1$  distance:  $\|\mathbf{a} - \mathbf{b}\|_1 = \sum_j |a_j - b_j|$ .
  - $\ell_2$  distance:  $\|\mathbf{a} - \mathbf{b}\|_2 = \sqrt{\sum_j (a_j - b_j)^2}$ .
  - $\ell_\infty$  distance:  $\|\mathbf{a} - \mathbf{b}\|_\infty = \max_j |a_j - b_j|$ .
  - $\ell_p$  distance:  $\|\mathbf{a} - \mathbf{b}\|_p = [\sum_j (a_j - b_j)^p]^{1/p}$ .

# Distance between Two Probability Distributions

**Question:** What is the distance between  $p(x)$  and  $q(x)$ ?

- Let  $p(x)$  and  $q(x)$  be two Probability Density Functions (PDFs).
- E.g., the PDF of a Gaussian distribution is  $p(x|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$ .



# Distance between Two Probability Distributions

**Question:** What is the distance between  $p(x)$  and  $q(x)$ ?

- Let  $p(x)$  and  $q(x)$  be two Probability Density Functions (PDFs).
- E.g., the PDF of a Gaussian distribution is  $p(x|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$ .

**Definition:** The Kullback–Leibler (KL) divergence  $D_{\text{KL}}(p||q)$ .

- For discrete probability distributions, the KL divergence is

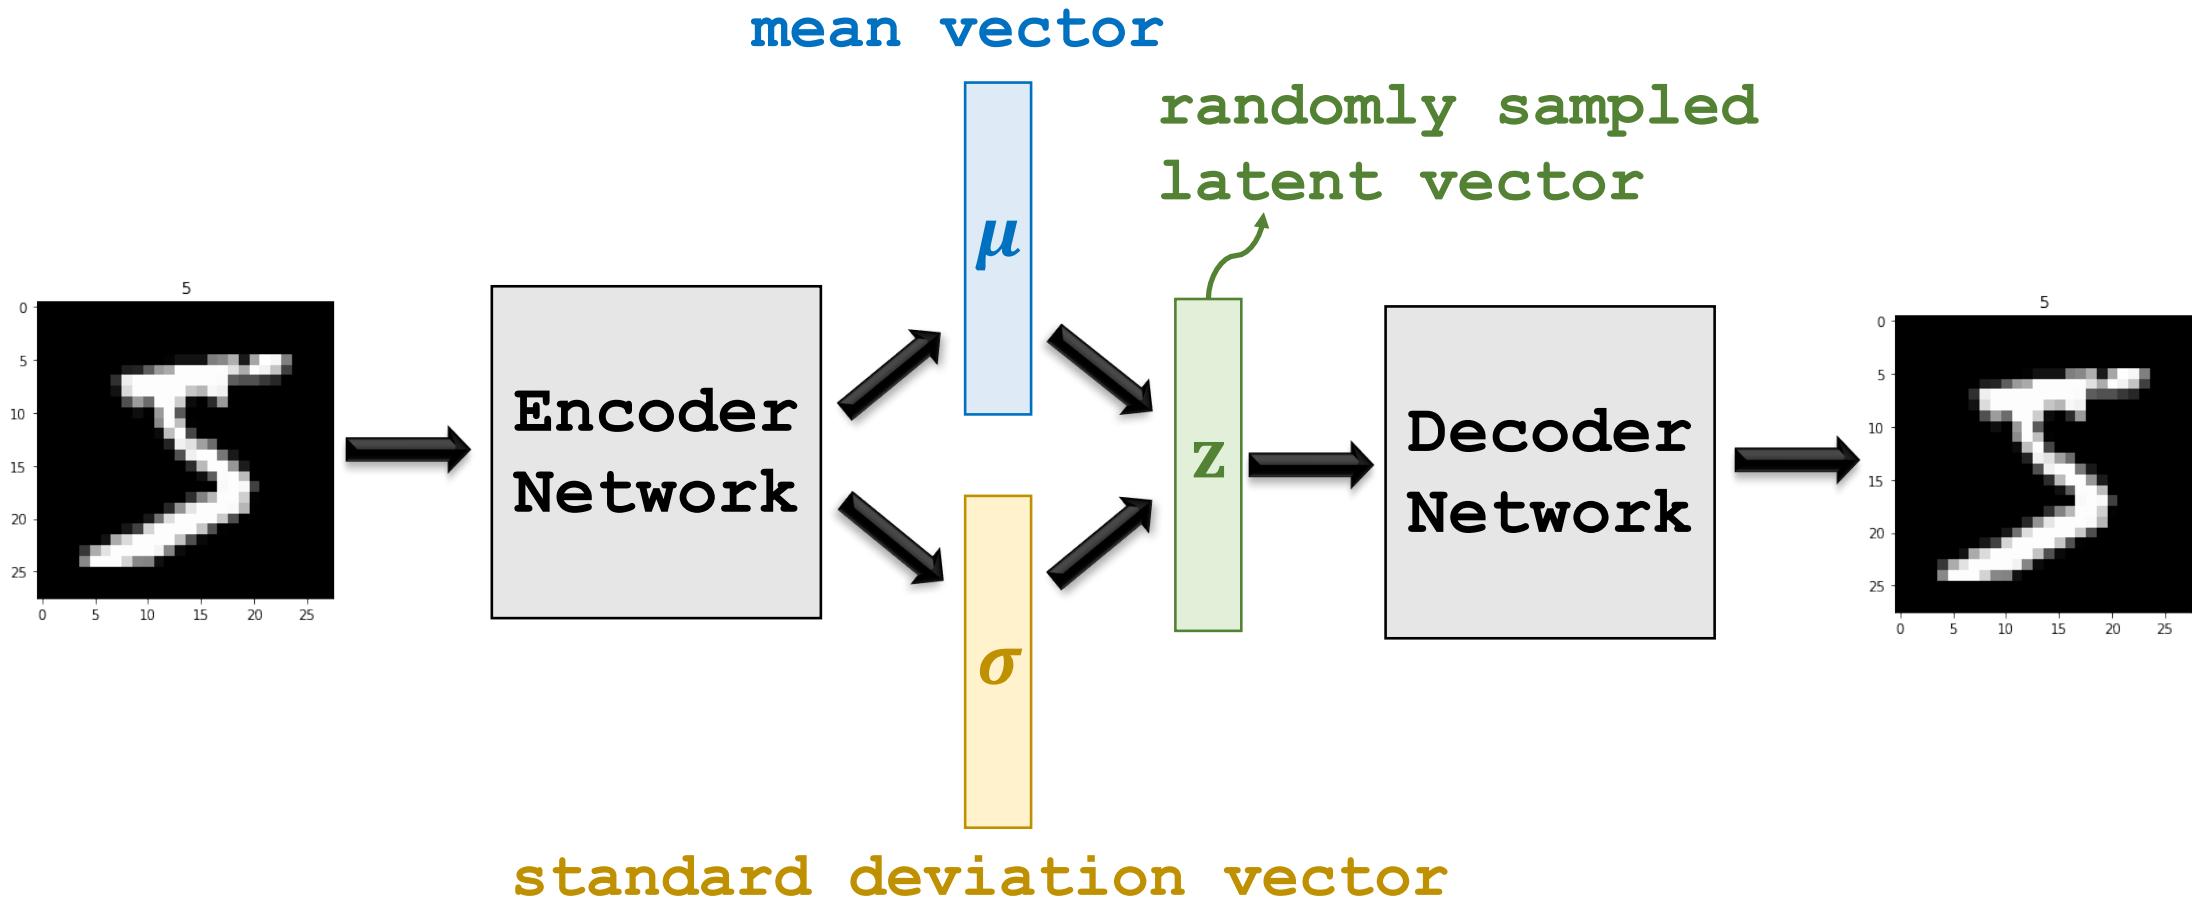
$$D_{\text{KL}}(p||q) = -\sum_i p(i) \log \frac{q(i)}{p(i)}.$$

- For continuous probability distributions, the KL divergence is

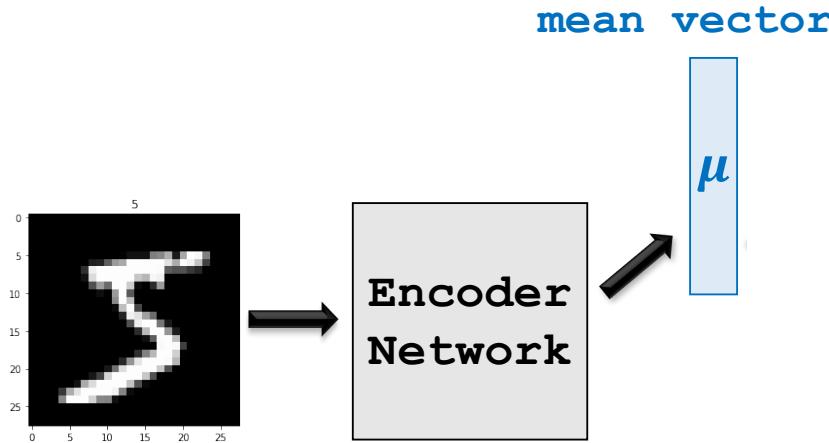
$$D_{\text{KL}}(p||q) = -\int_{-\infty}^{+\infty} p(x) \log \frac{q(x)}{p(x)} dx.$$

# Variational Autoencoder

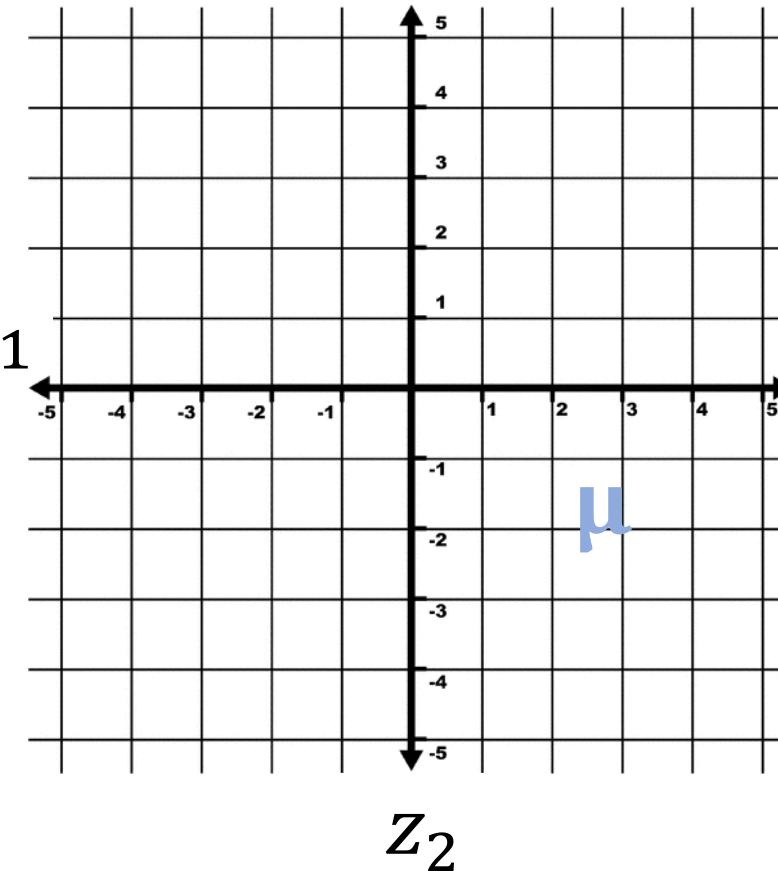
# Variational Autoencoder (VAE)



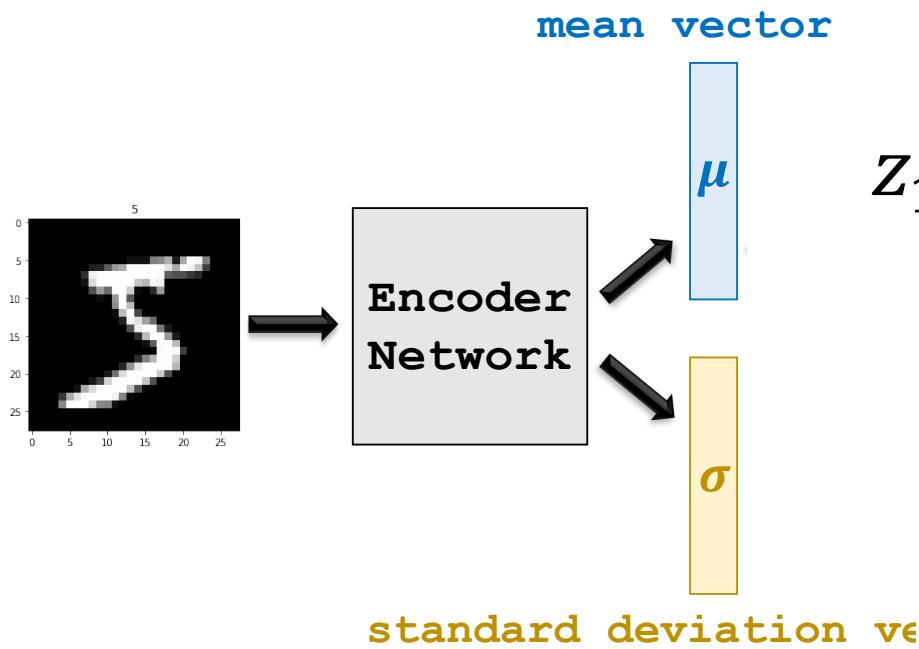
# VAE Sampling



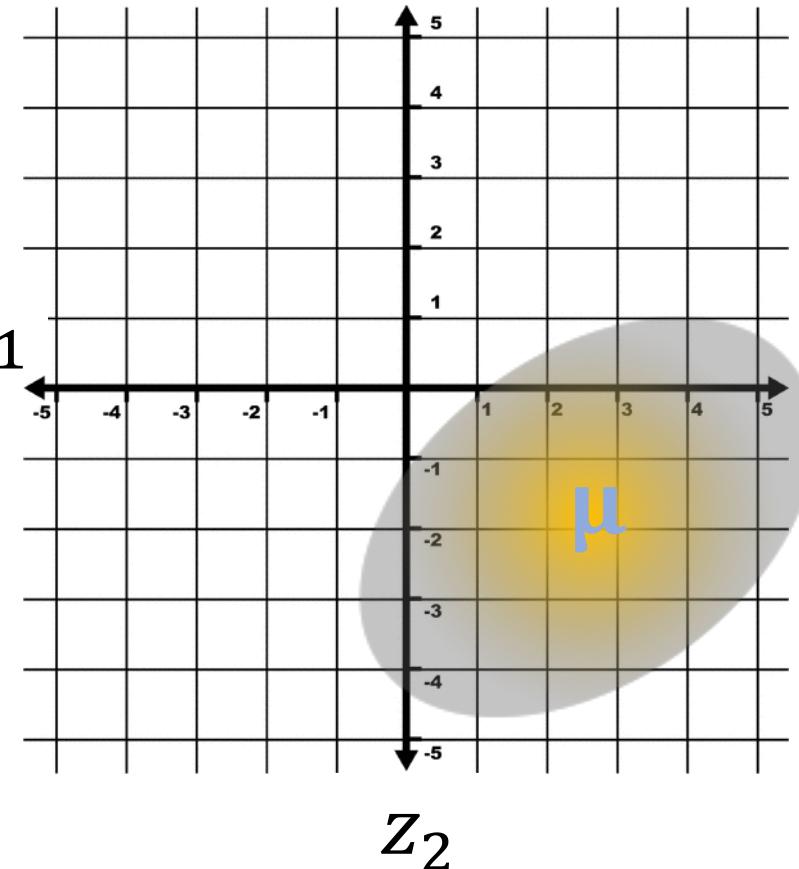
The latent space



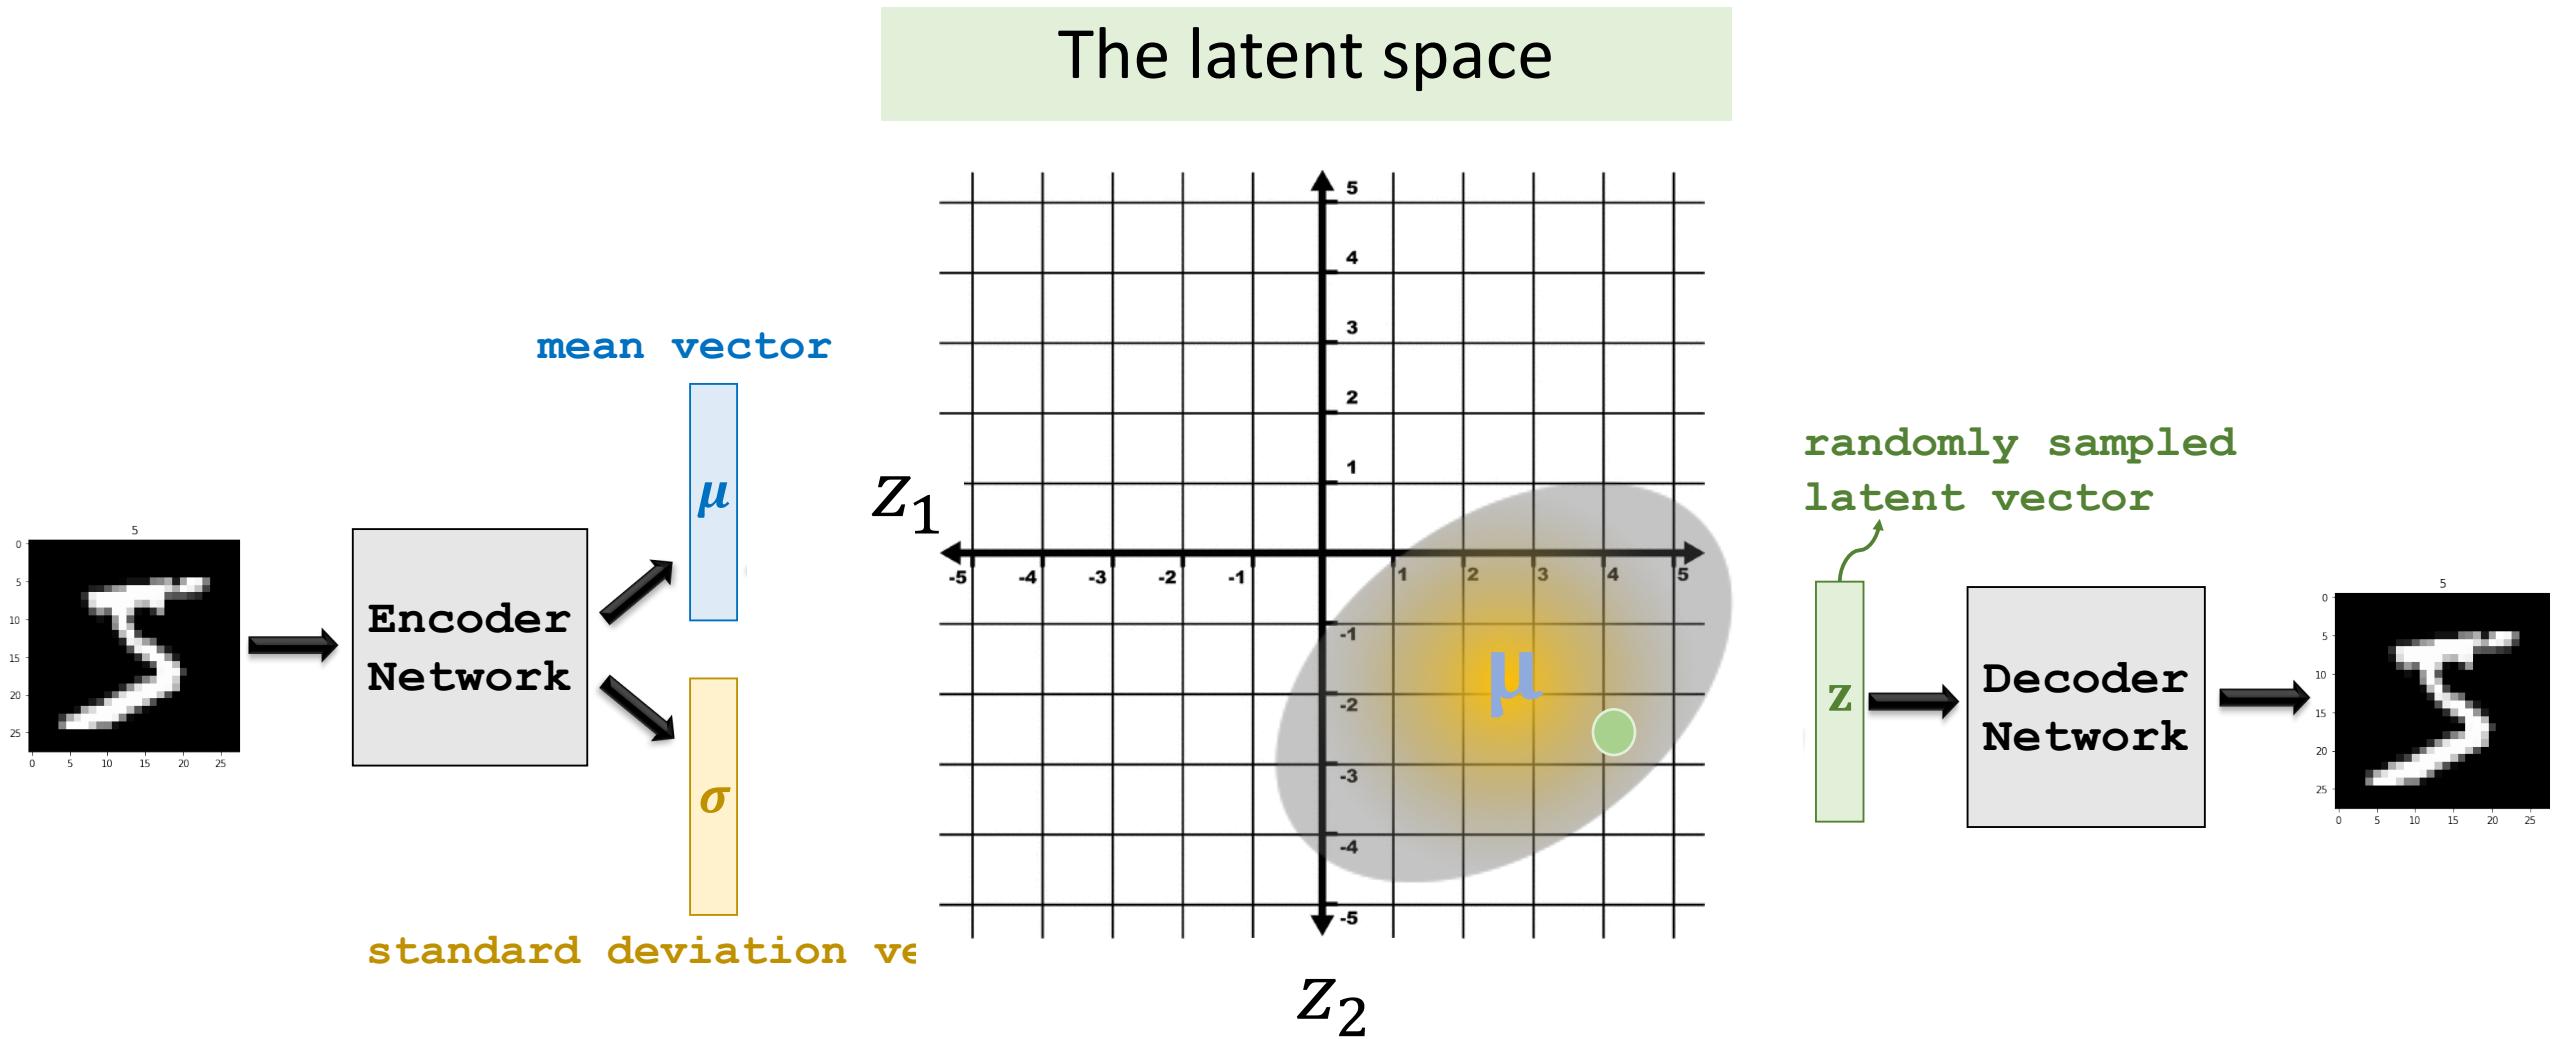
# VAE Sampling



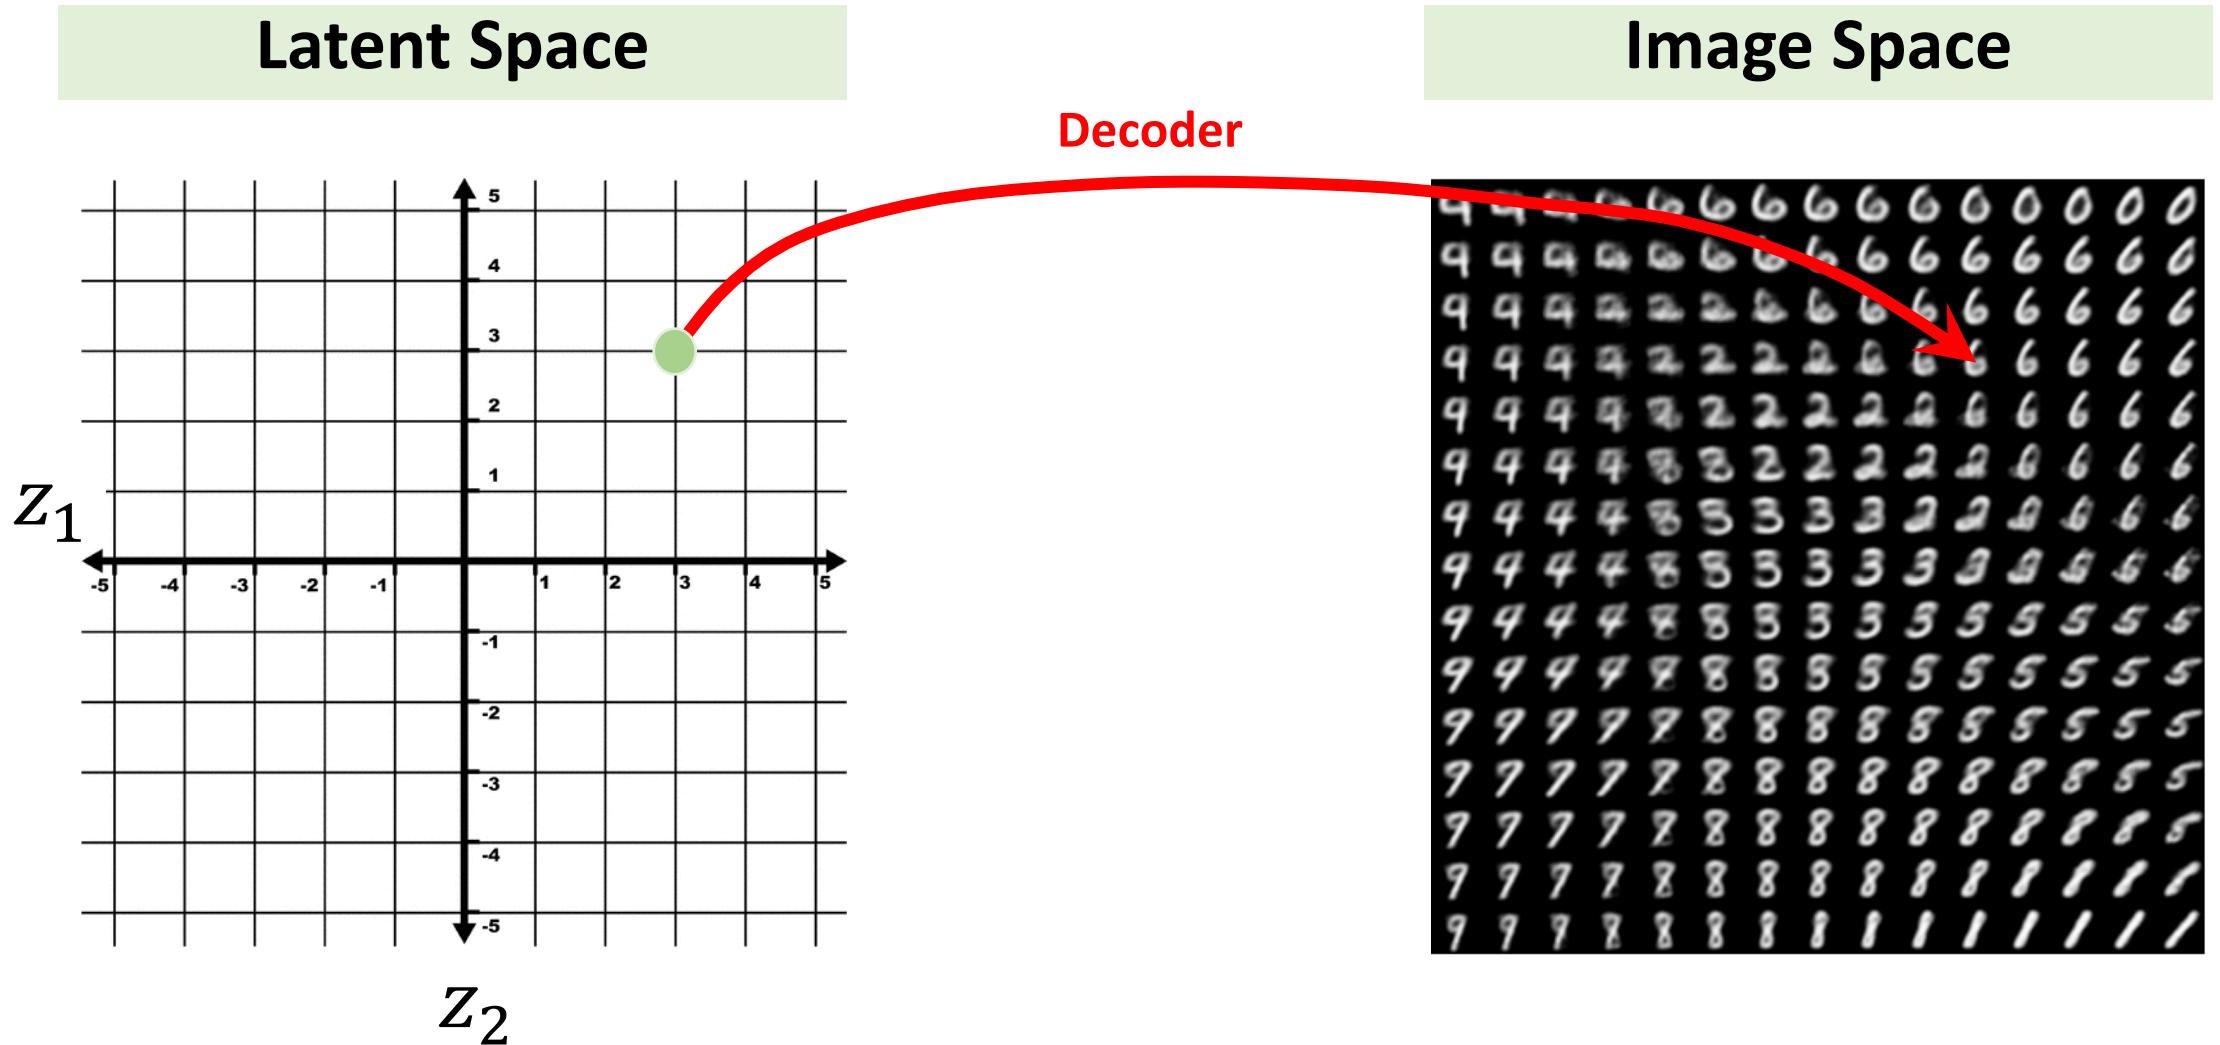
The latent space



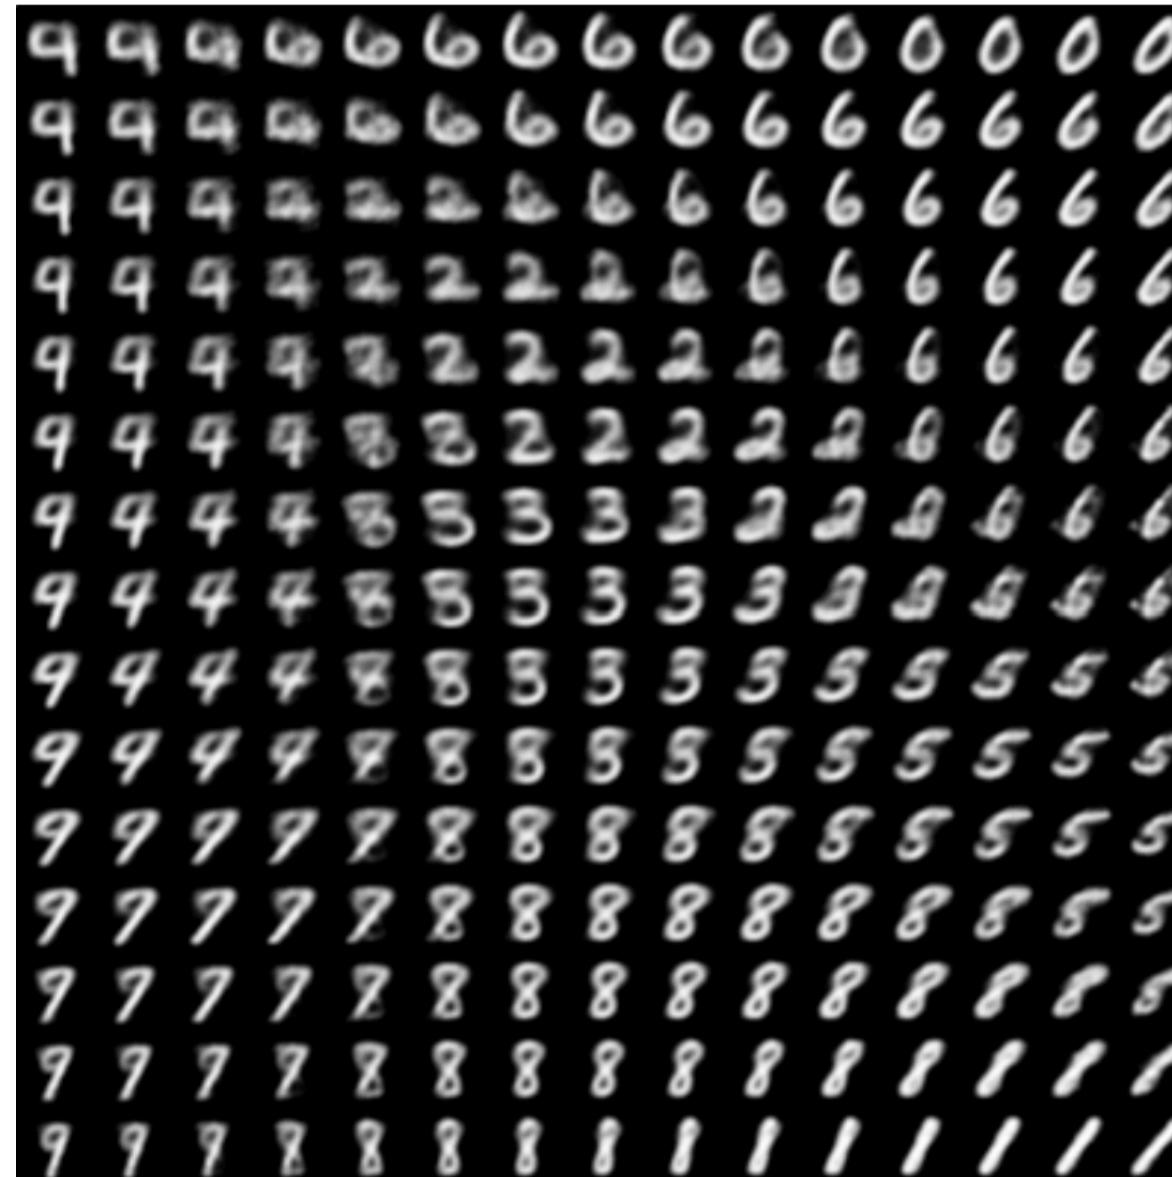
# VAE Sampling



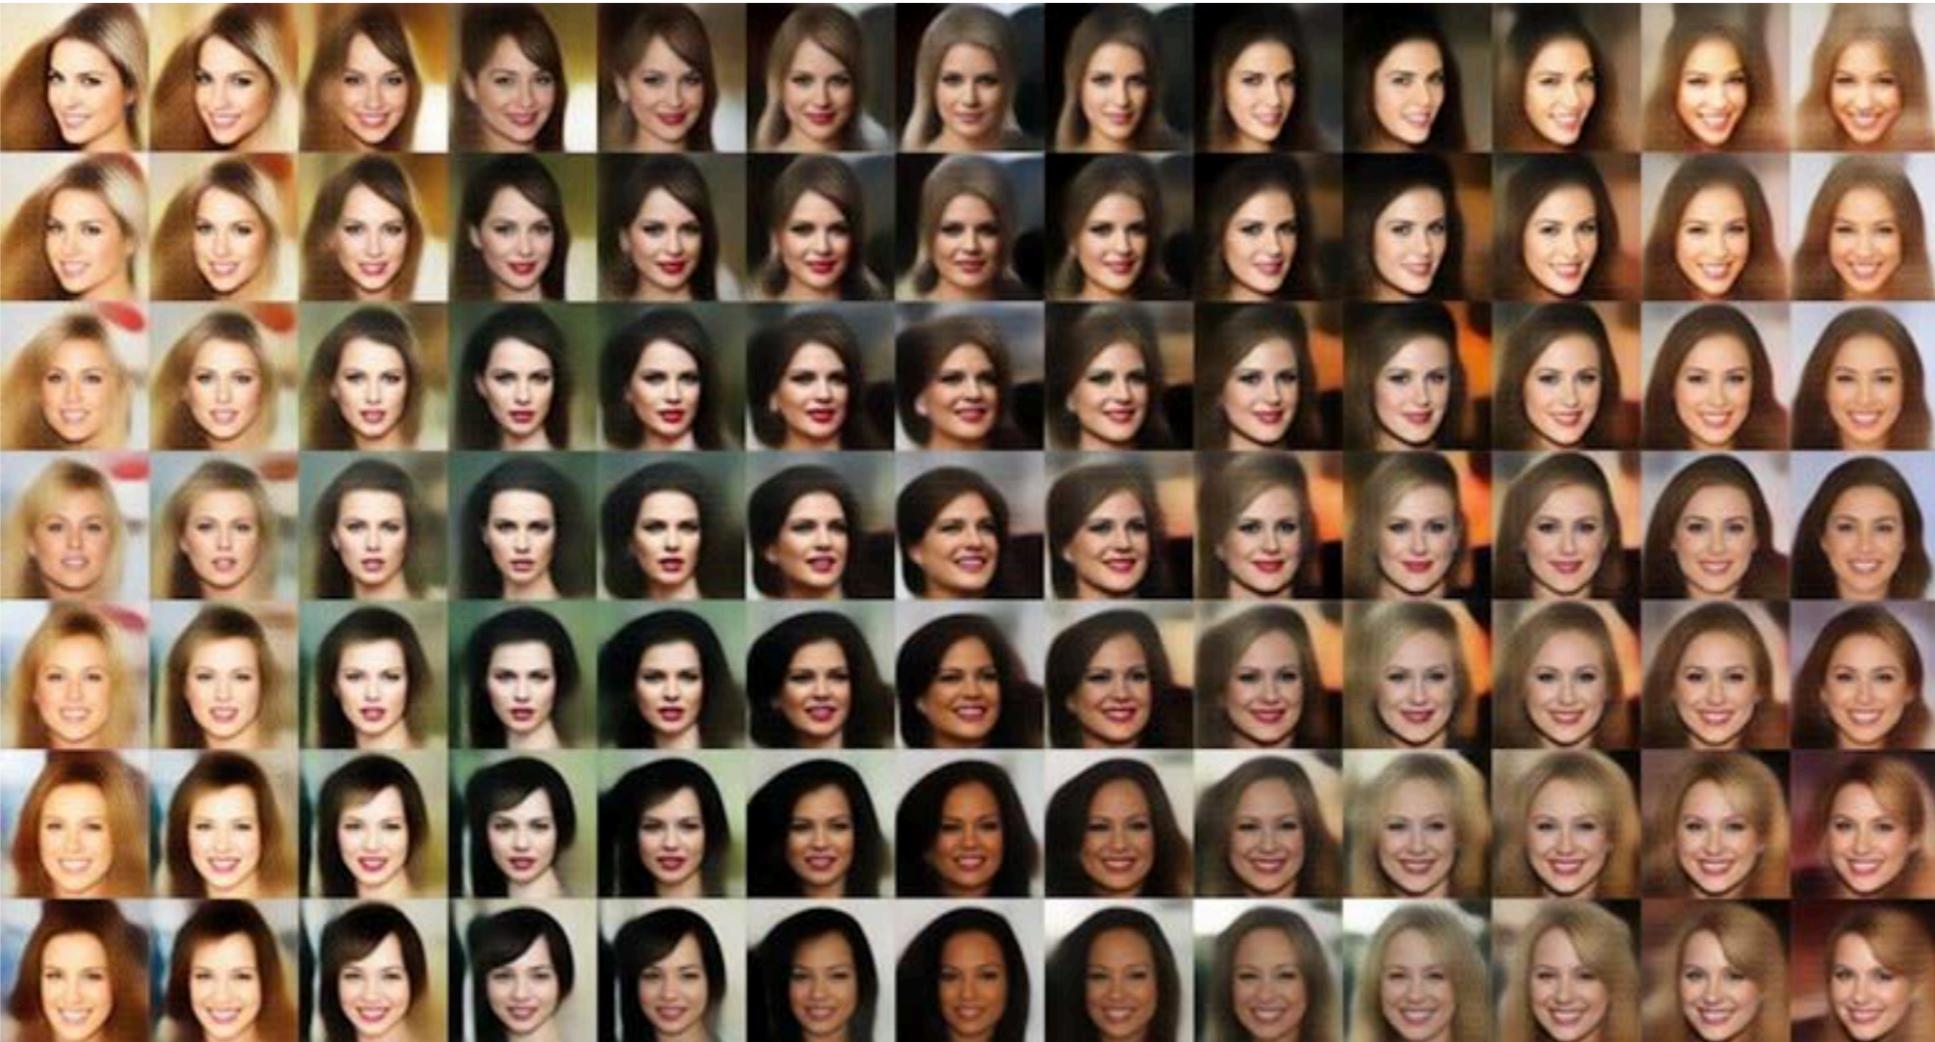
# Visualize the Latent Space



# The Learned Latent Space is Continuous!

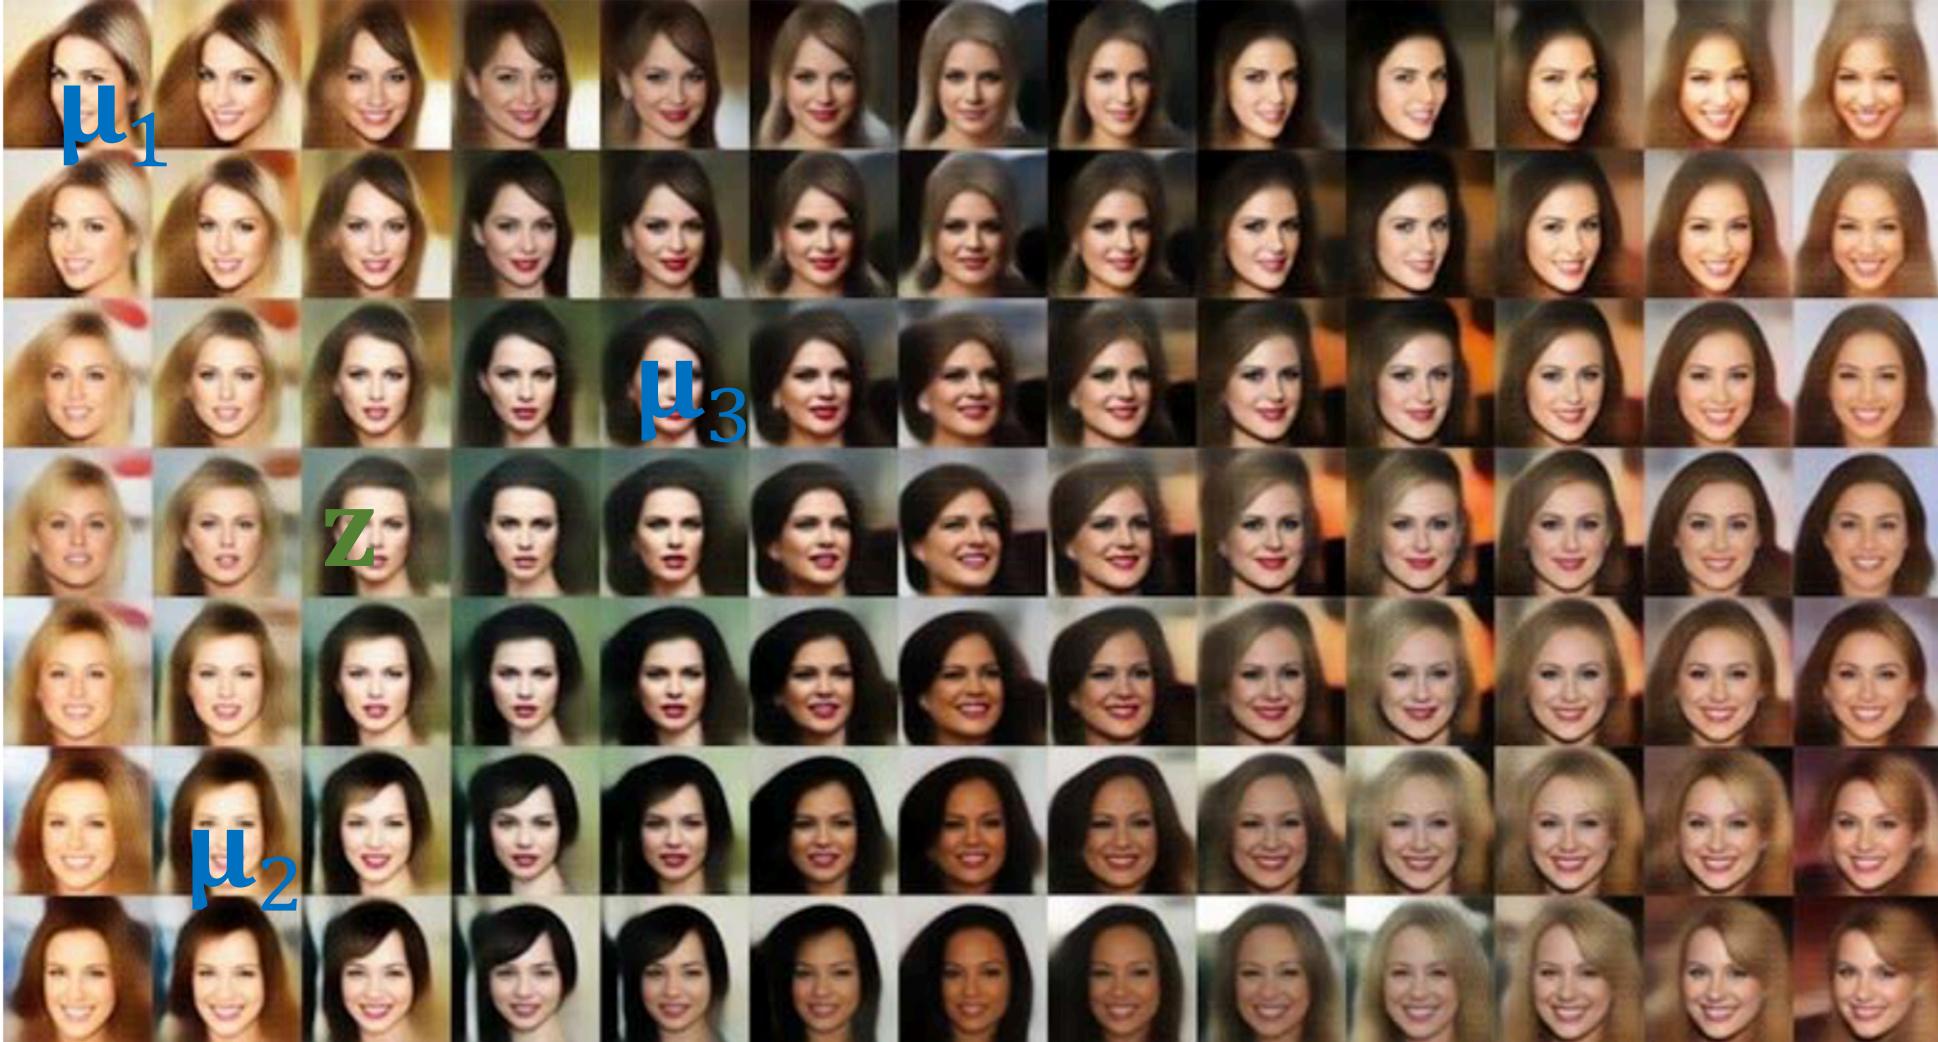


# The Learned Latent Space is Continuous!



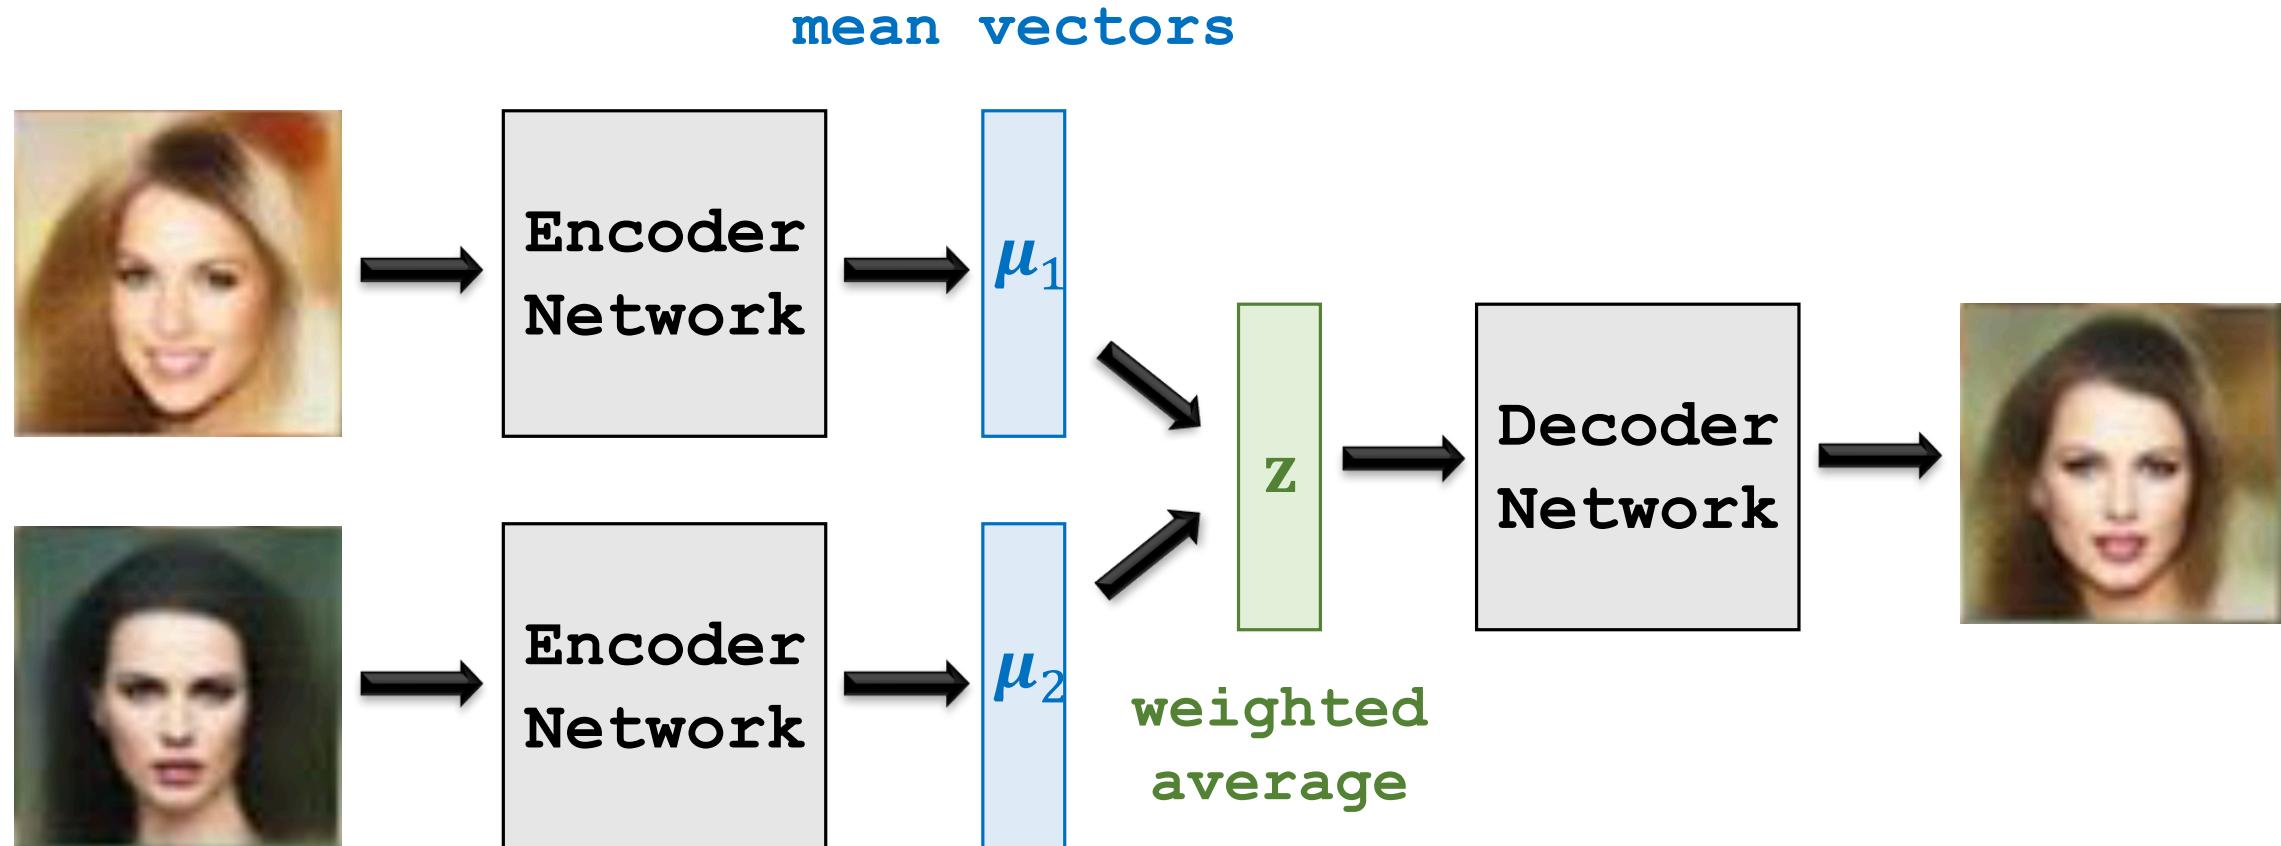
Visualization of the latent space of faces (by Tom White)

# Average Images in the Latent Space

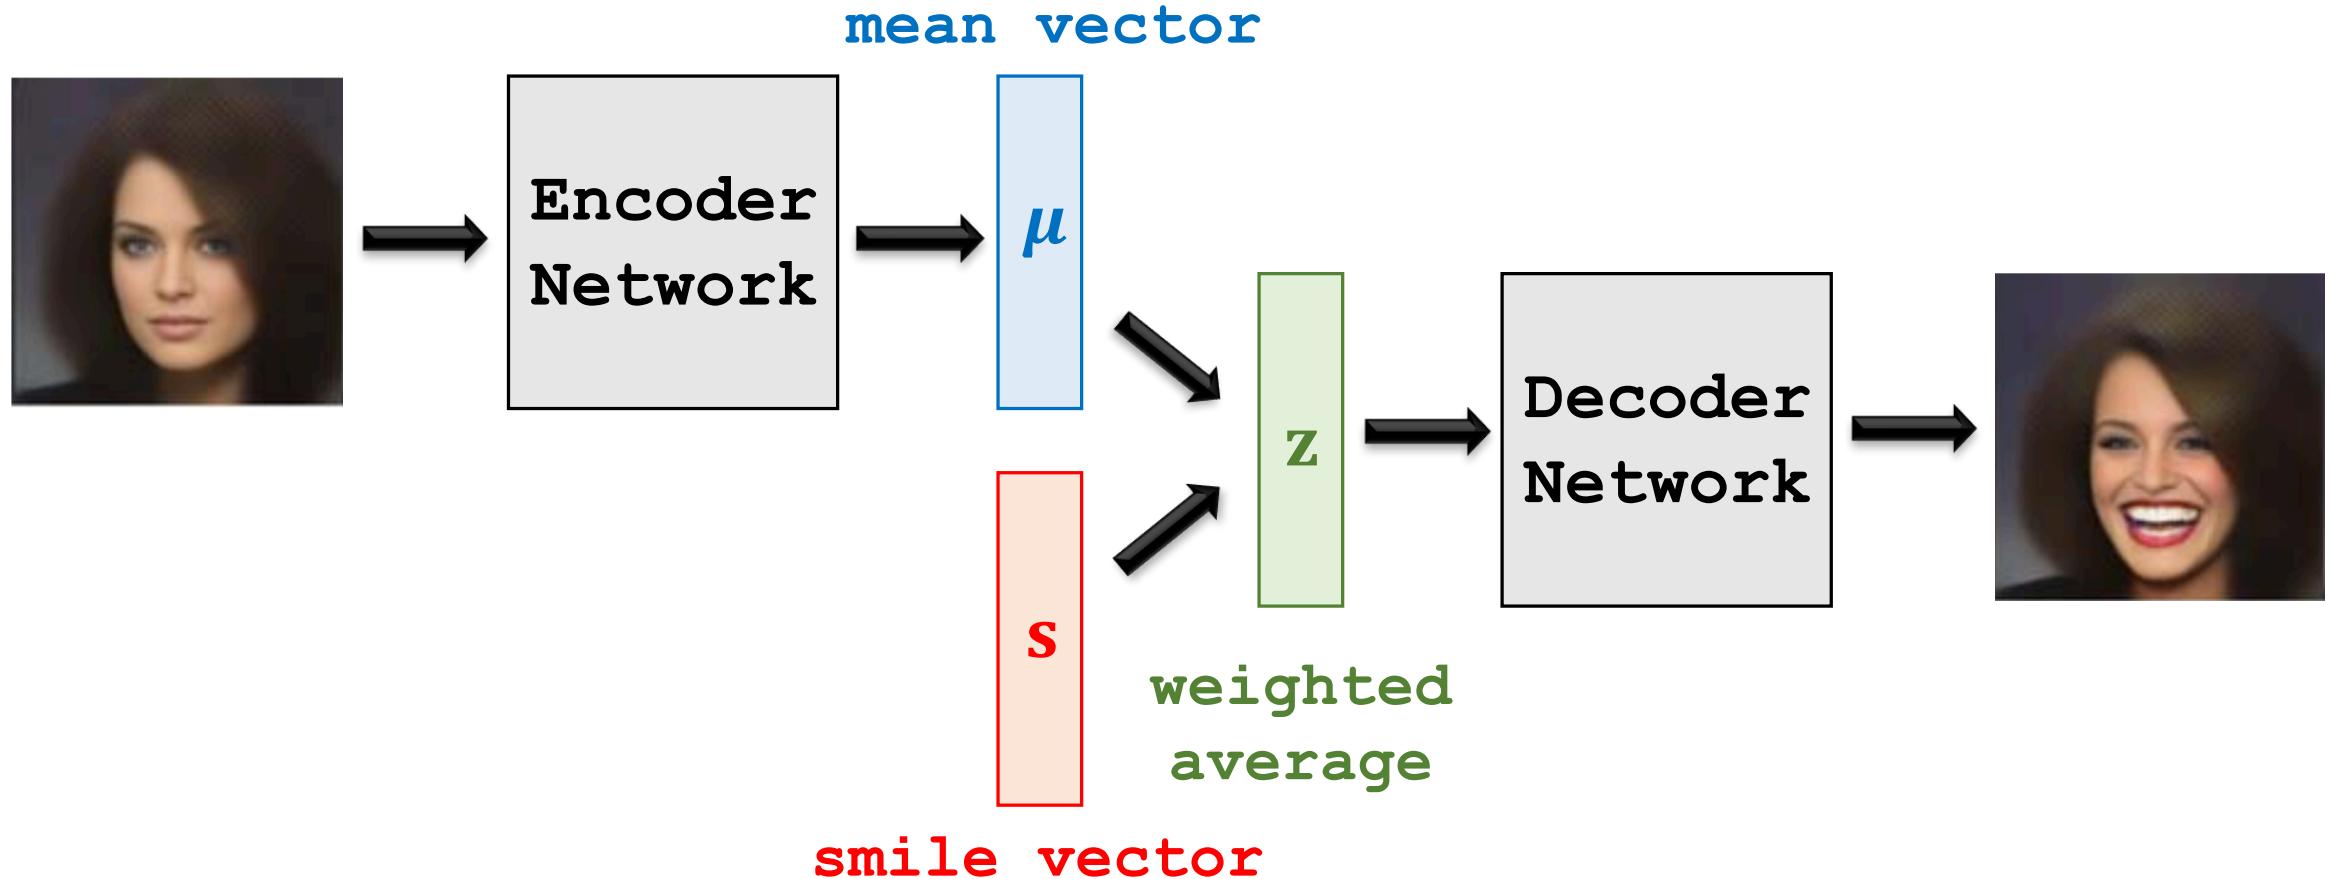


Visualization of the latent space of faces (by Tom White)

# Average Images in the Latent Space



# Edit Images in the Latent Space



# Edit Images in the Latent Space

$\mu$

$\mu + s$

$\mu + 2s$

$\mu + 3s$

$\mu + 4s$



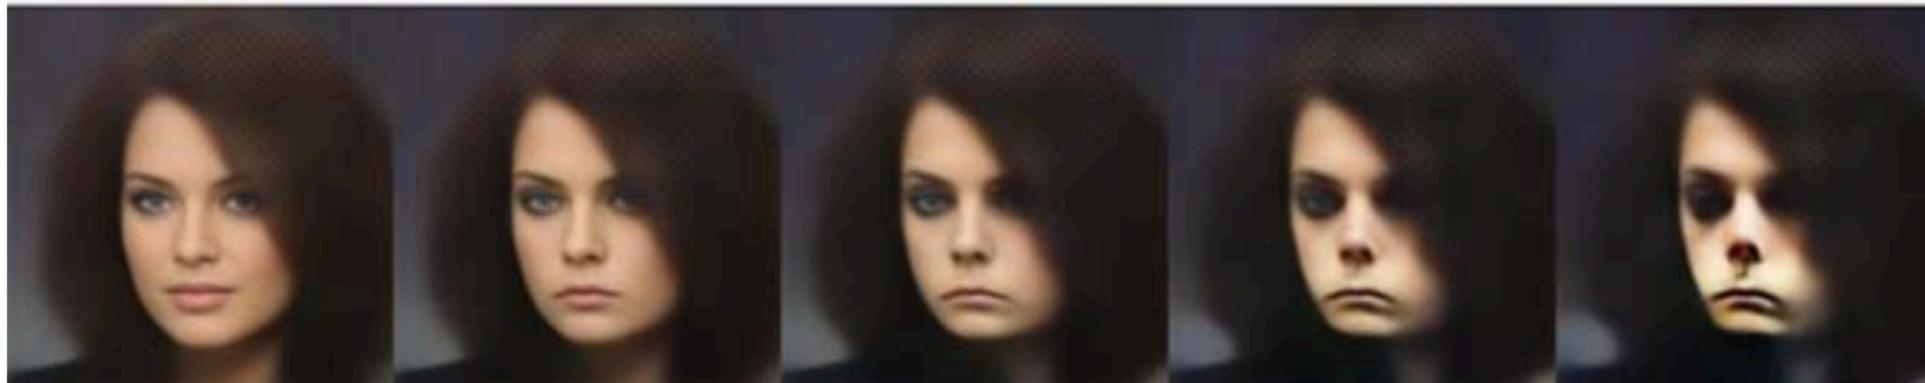
$\mu$

$\mu - s$

$\mu - 2s$

$\mu - 3s$

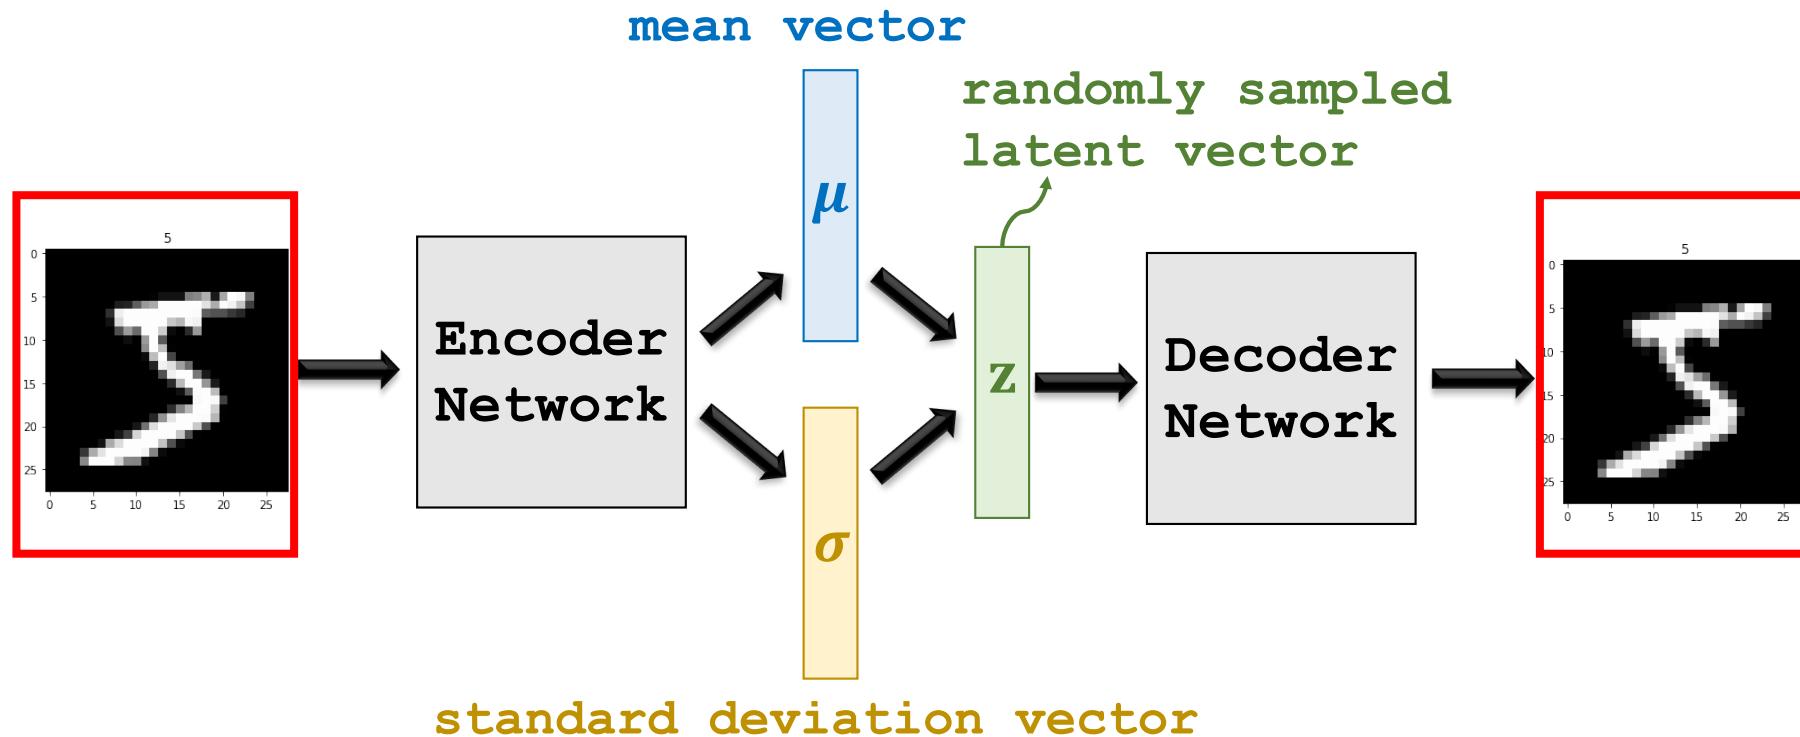
$\mu - 4s$



# What makes the Latent Space Continuous?

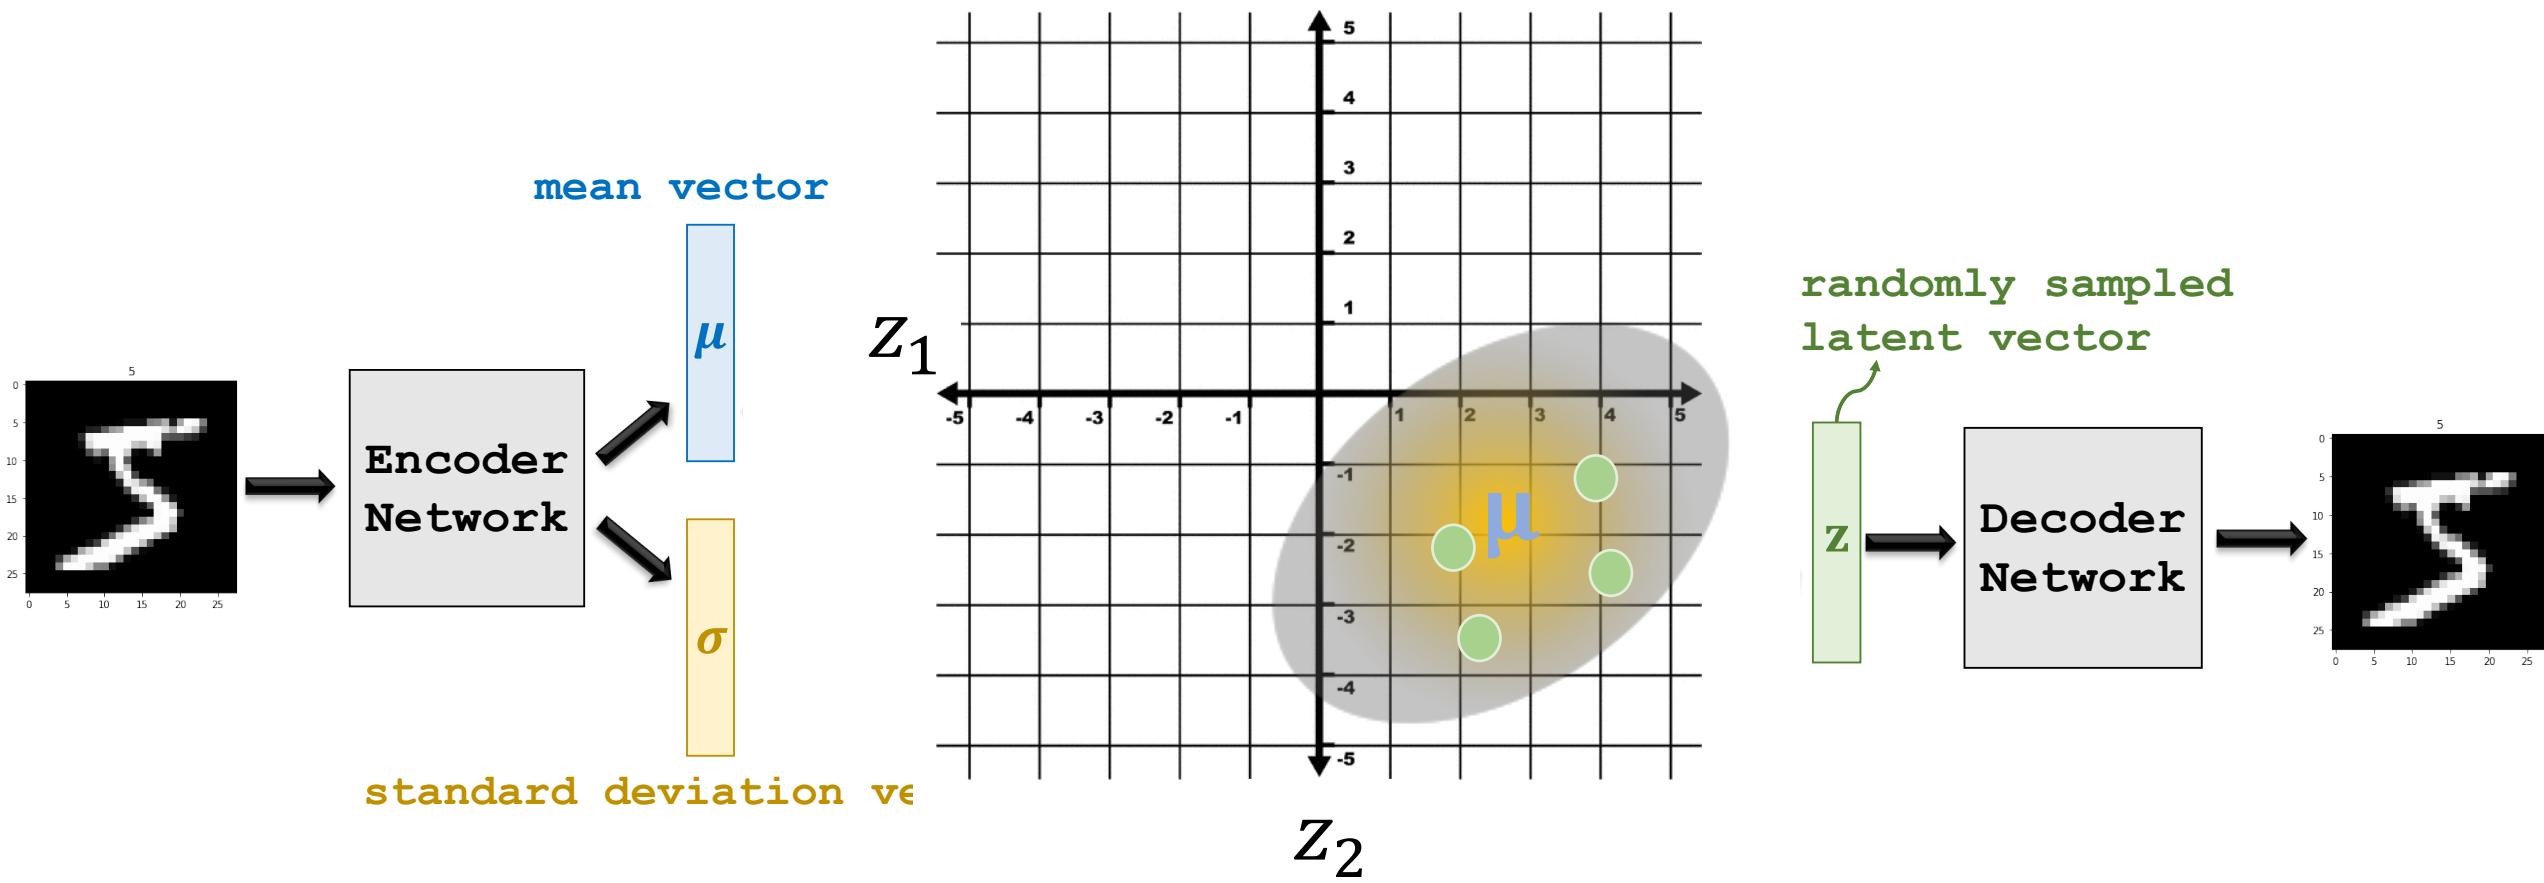
# The Generation Loss

- $\text{GenLoss} = \text{dist}(\text{input\_img}, \text{generated\_img})$ .
- E.g., the  $\ell_2$  distance, the cross-entropy, etc.



# The Generation Loss

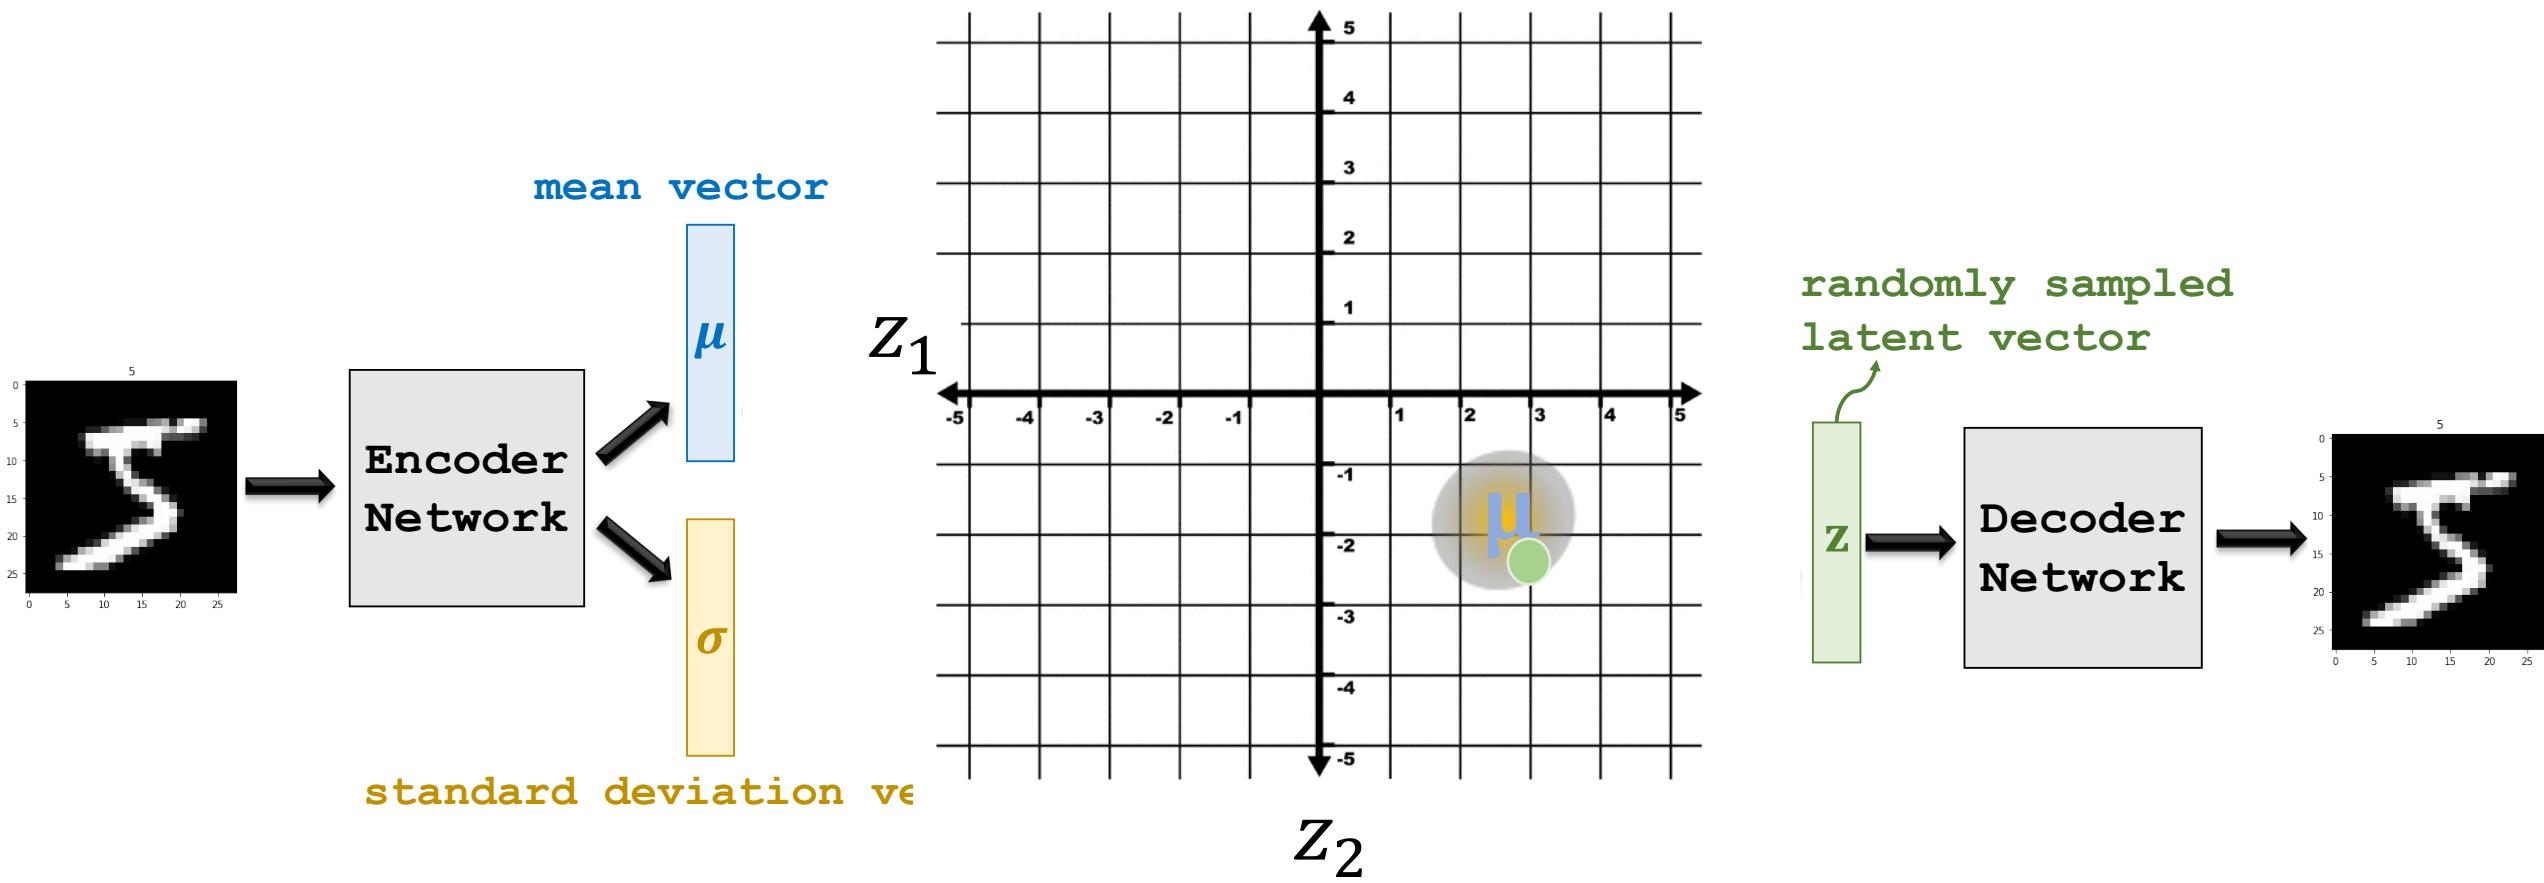
- In the latent space,  $\mathbf{z}$  is close to  $\mu$ , but  $\mathbf{z} \neq \mu$ .
- All the latent vectors around  $\mu$  lead to similar generated images.
- Thus the latent space is “*continuous*” (in terms of the generated images).



# The Generation Loss

**Difficulty:** The encoder network will learn an std vector  $\sigma \rightarrow 0$ . (Why?)

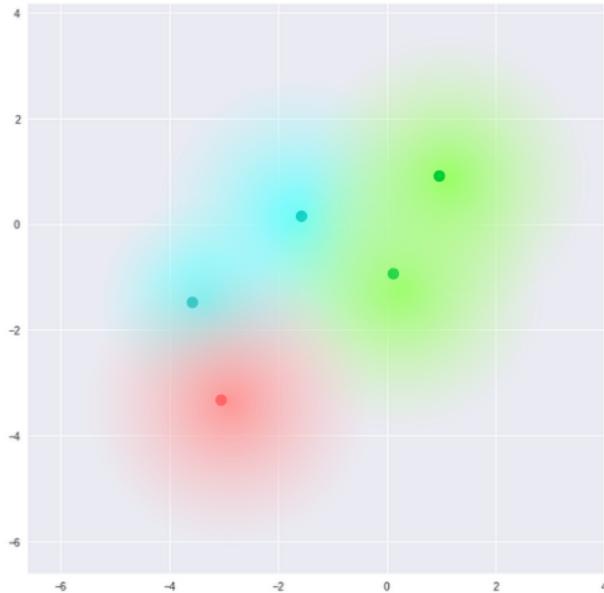
- Minimize GenLoss  $\rightarrow$  encourage  $\mathbf{z}$  close to  $\boldsymbol{\mu}$   $\rightarrow$  encourage small  $\sigma$ .
- VAE degrades to the standard AE  $\rightarrow$  the latent space is not continuous.



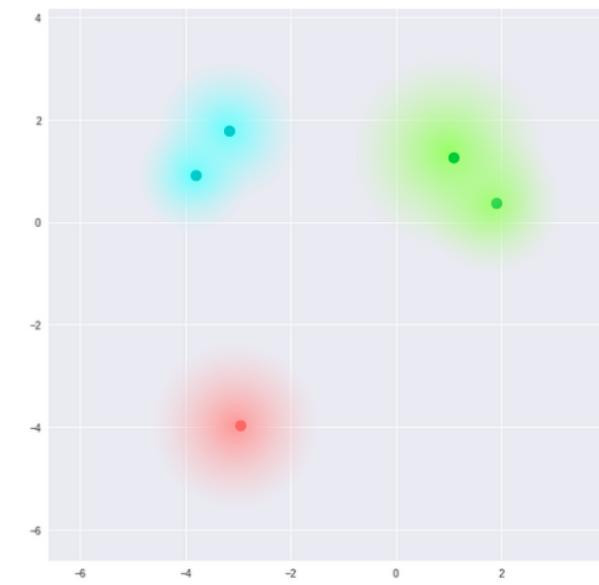
# The Generation Loss

**Difficulty:** The encoder network will learn an std vector  $\sigma \rightarrow \mathbf{0}$ . (Why?)

- Minimize GenLoss  $\rightarrow$  encourage  $\mathbf{z}$  close to  $\mu$   $\rightarrow$  encourage small  $\sigma$ .
- VAE degrades to the standard AE  $\rightarrow$  the latent space is not continuous.



The continuous latent space which we hope to learn.

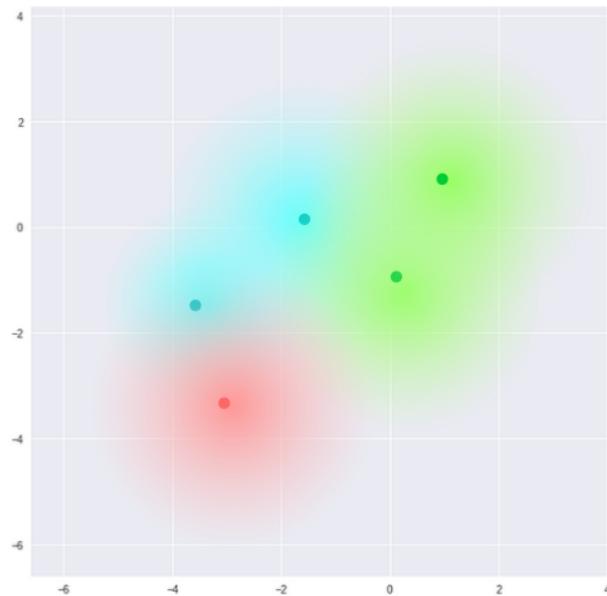


What we actually learn by minimizing GenLoss.

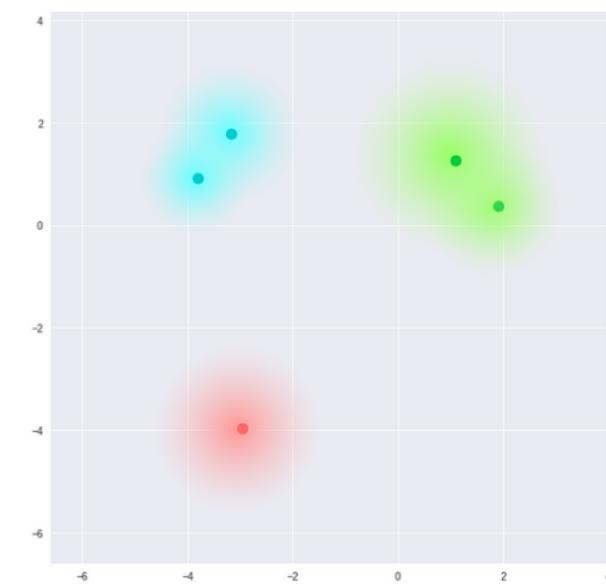
# The Latent Loss

Idea: Encourage the distribution of  $\mathbf{z}$  to be unit Gaussian.

- $\text{LatLoss} = D_{\text{KL}}(p(\mathbf{z}) \parallel \mathcal{N}(0, \mathbf{I}))$



The latent space which we hope to learn.



What we actually learn by minimizing GenLoss.

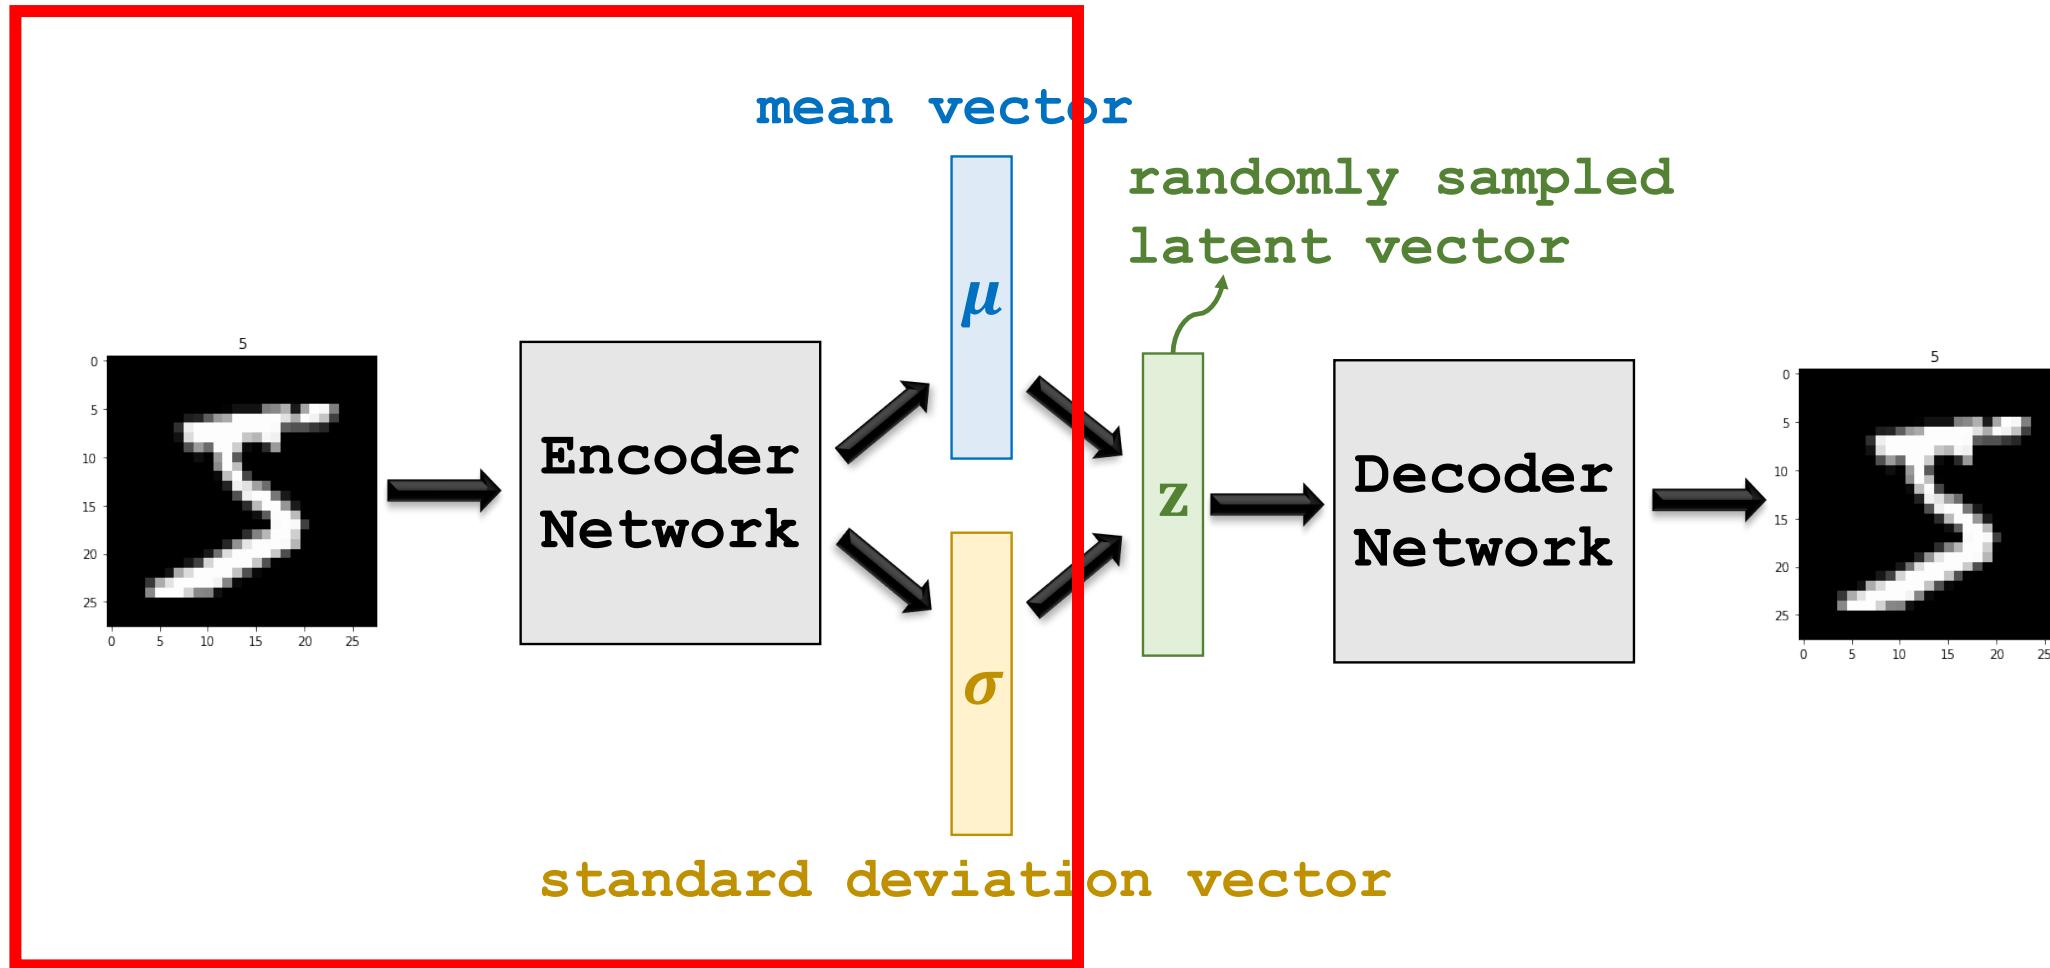
# The Latent Loss

Idea: Encourage the distribution of  $\mathbf{z}$  to be unit Gaussian.

- $\text{LatLoss} = D_{\text{KL}}(p(\mathbf{z}) \parallel \mathcal{N}(0, \mathbf{I}))$
- $\text{Loss} = \text{GenLoss} + \lambda \cdot \text{LatLoss}$

# Build the Encoder

# 1. The Encoder Network



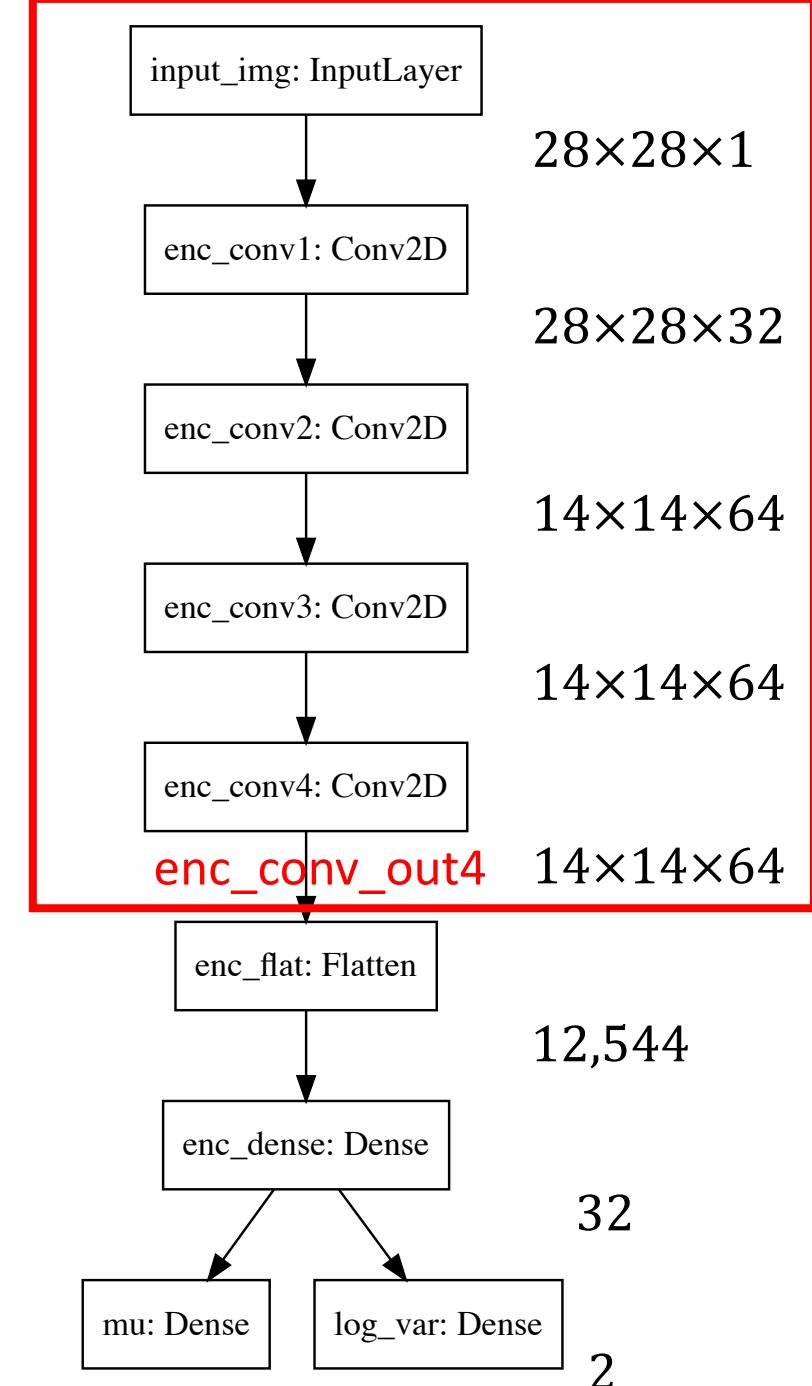
# 1. The Encoder Network

```
from keras.layers import Input, Conv2D, Flatten, Dense
from keras import models

latent_dim = 2

# define convolutional layers
enc_conv1 = Conv2D(32, 3, padding='same',
                  activation='relu', name='enc_conv1')
enc_conv2 = Conv2D(64, 3, padding='same', activation='relu',
                  strides=(2, 2), name='enc_conv2')
enc_conv3 = Conv2D(64, 3, padding='same',
                  activation='relu', name='enc_conv3')
enc_conv4 = Conv2D(64, 3, padding='same',
                  activation='relu', name='enc_conv4')

input_img = Input(shape=(28, 28, 1), name='input_img')
enc_conv_out1 = enc_conv1(input_img)
enc_conv_out2 = enc_conv2(enc_conv_out1)
enc_conv_out3 = enc_conv3(enc_conv_out2)
enc_conv_out4 = enc_conv4(enc_conv_out2)
```

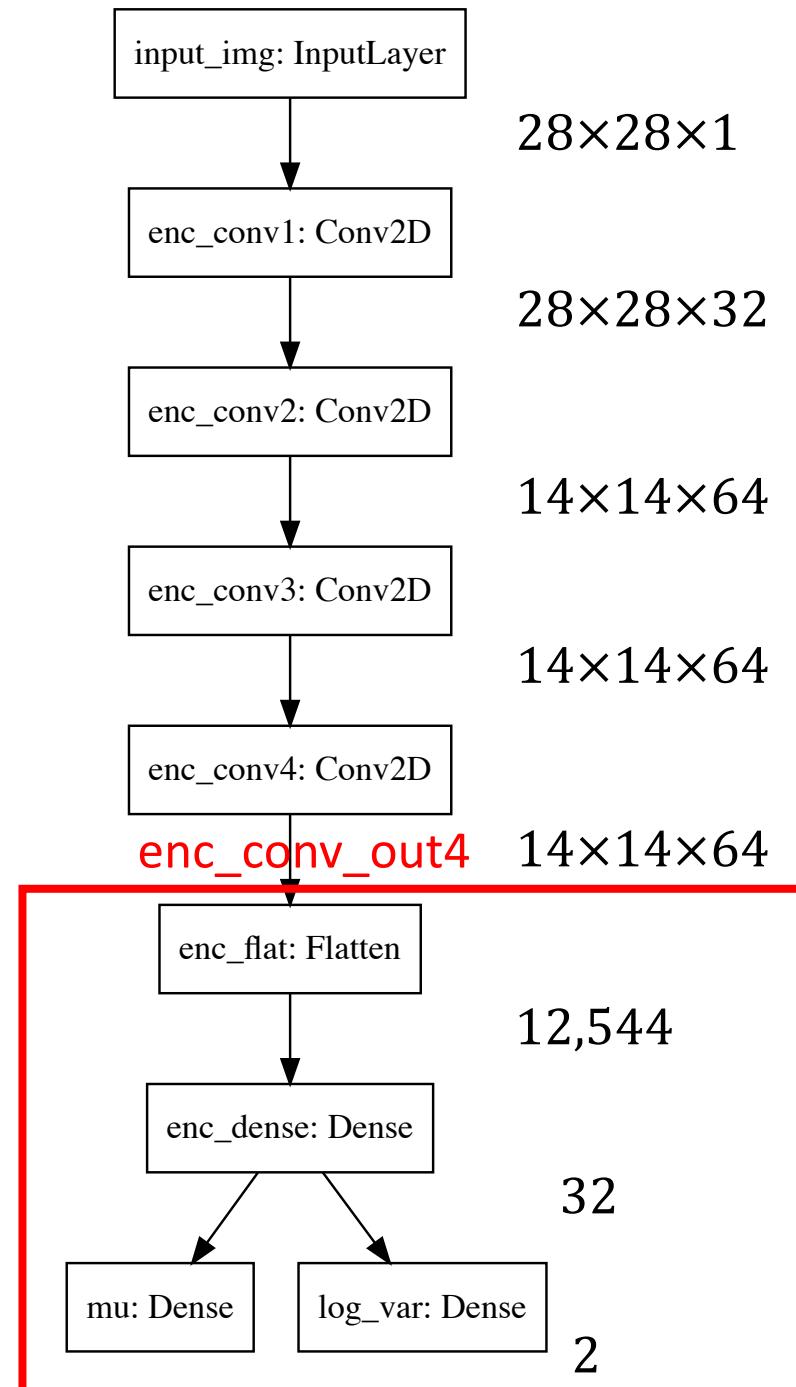


# 1. The Encoder Network

```
# define flatten and dense layers
enc_flat = Flatten(name='enc_flat')
enc_dense = Dense(32, activation='relu',
                  name='enc_dense')
enc_mu = Dense(latent_dim, name='mu')
enc_log_var = Dense(latent_dim, name='log_var')

enc_flat_out = enc_flat(enc_conv_out4)
enc_dense_out = enc_dense(enc_flat_out)
mu = enc_mu(enc_dense_out)
log_var = enc_log_var(enc_dense_out)
```

```
# model
encoder = models.Model(inputs=input_img,
                       outputs=[mu, log_var],
                       name='encoder')
```

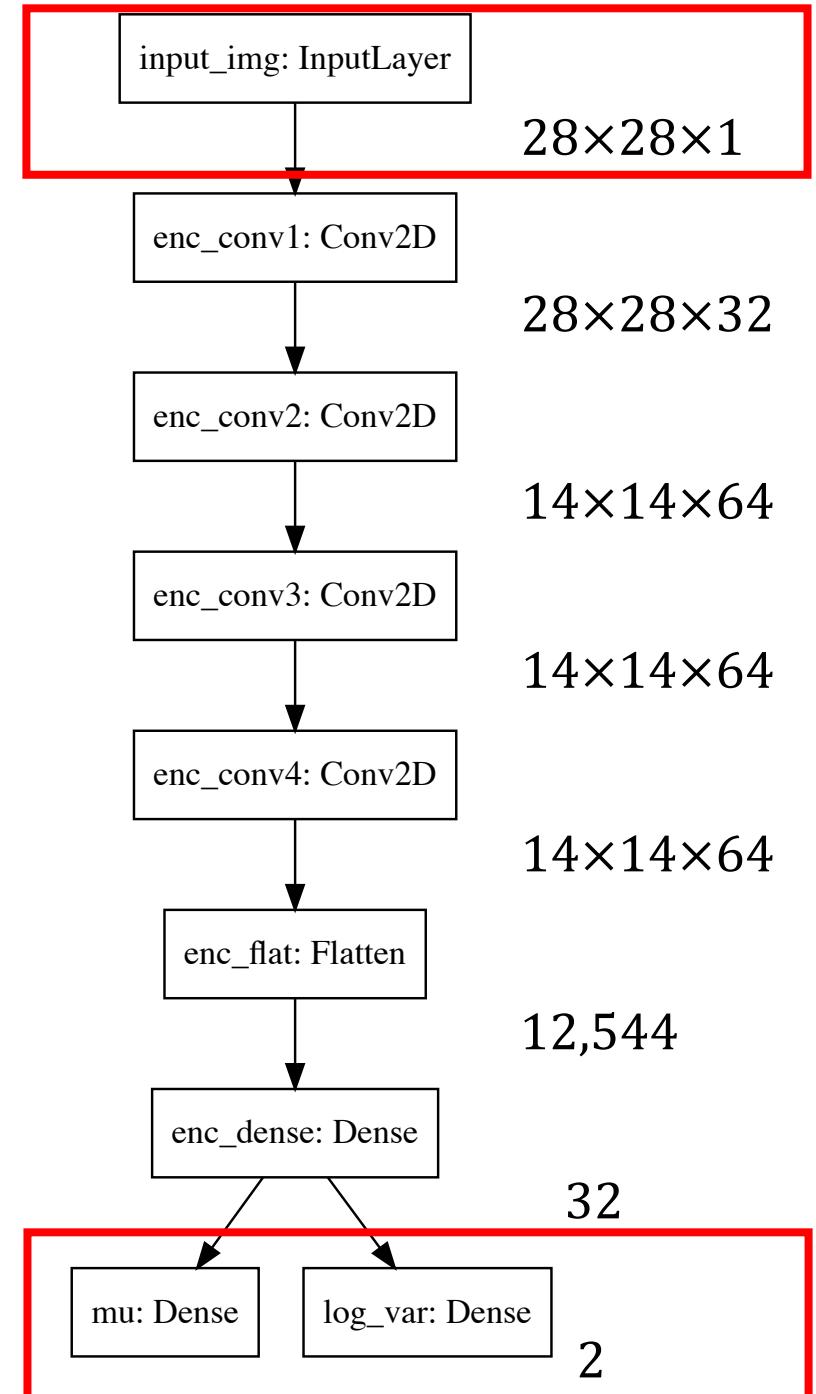


# 1. The Encoder Network

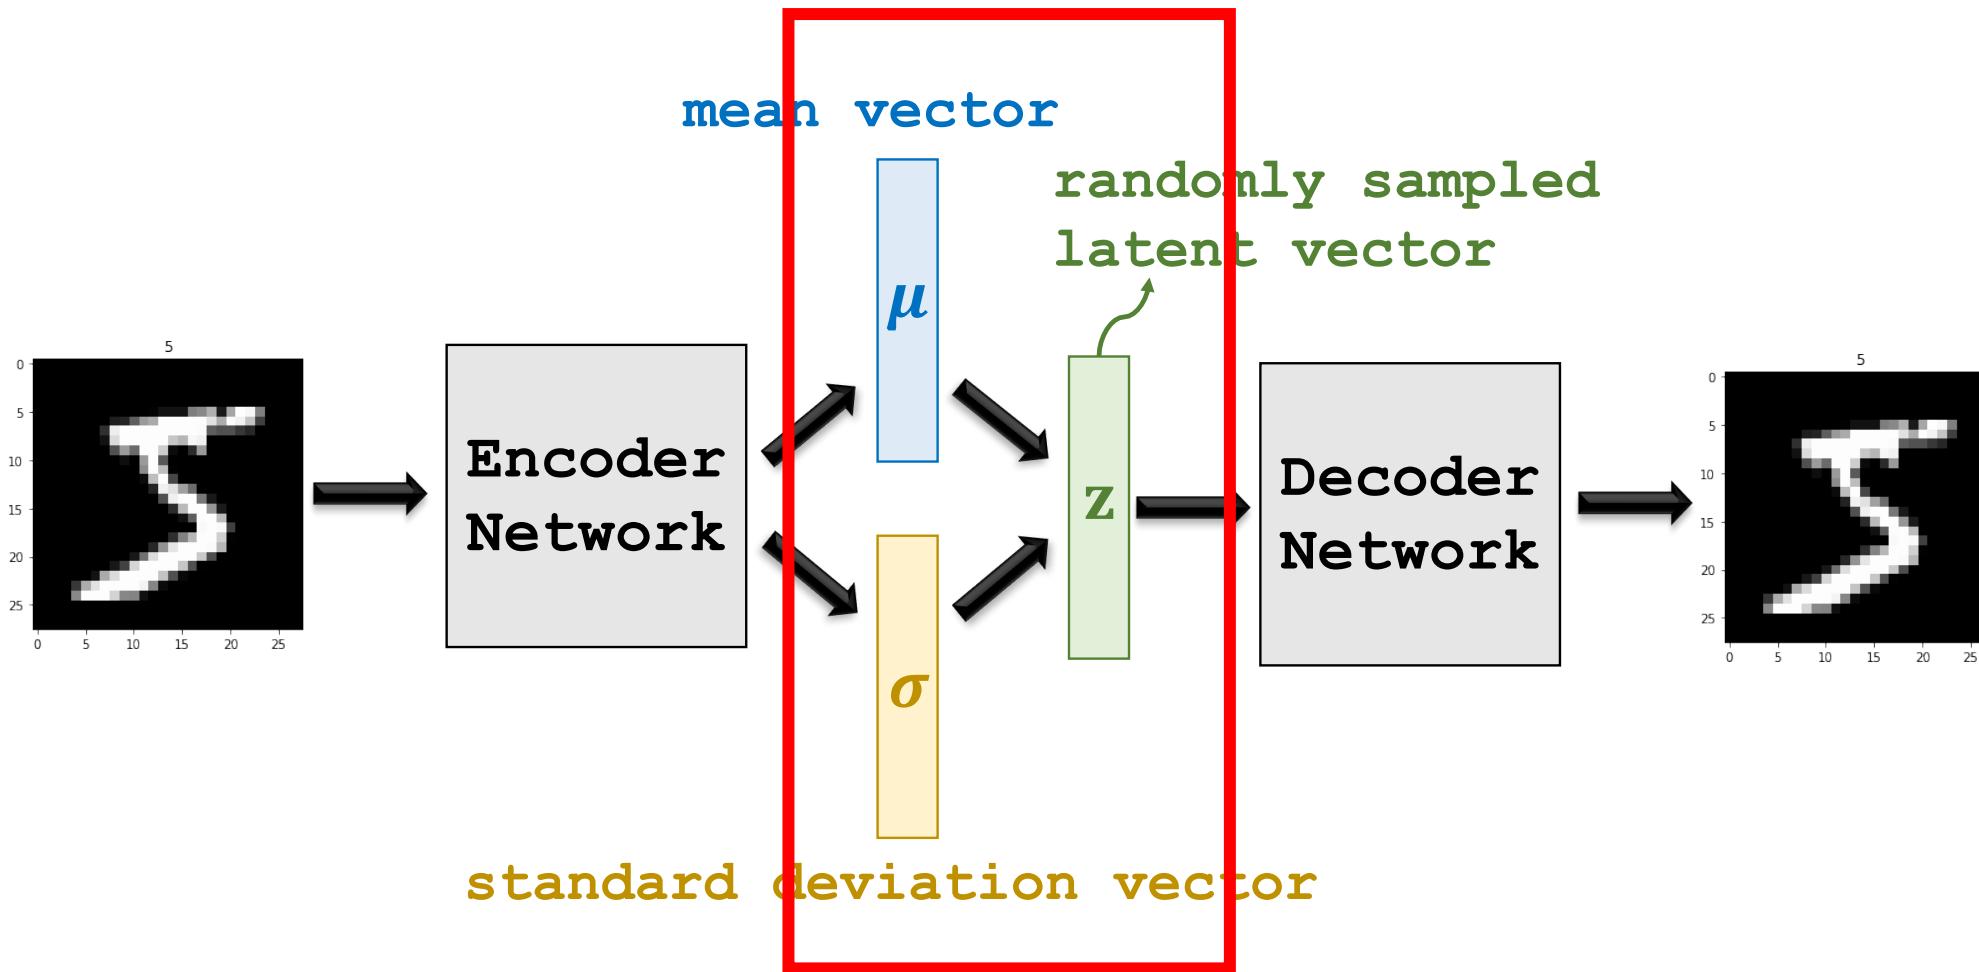
```
# define flatten and dense layers
enc_flat = Flatten(name='enc_flat')
enc_dense = Dense(32, activation='relu',
                  name='enc_dense')
enc_mu = Dense(latent_dim, name='mu')
enc_log_var = Dense(latent_dim, name='log_var')

enc_flat_out = enc_flat(enc_conv_out4)
enc_dense_out = enc_dense(enc_flat_out)
mu = enc_mu(enc_dense_out)
log_var = enc_log_var(enc_dense_out)

# model
encoder = models.Model(inputs=input_img,
                       outputs=[mu, log_var],
                       name='encoder')
```

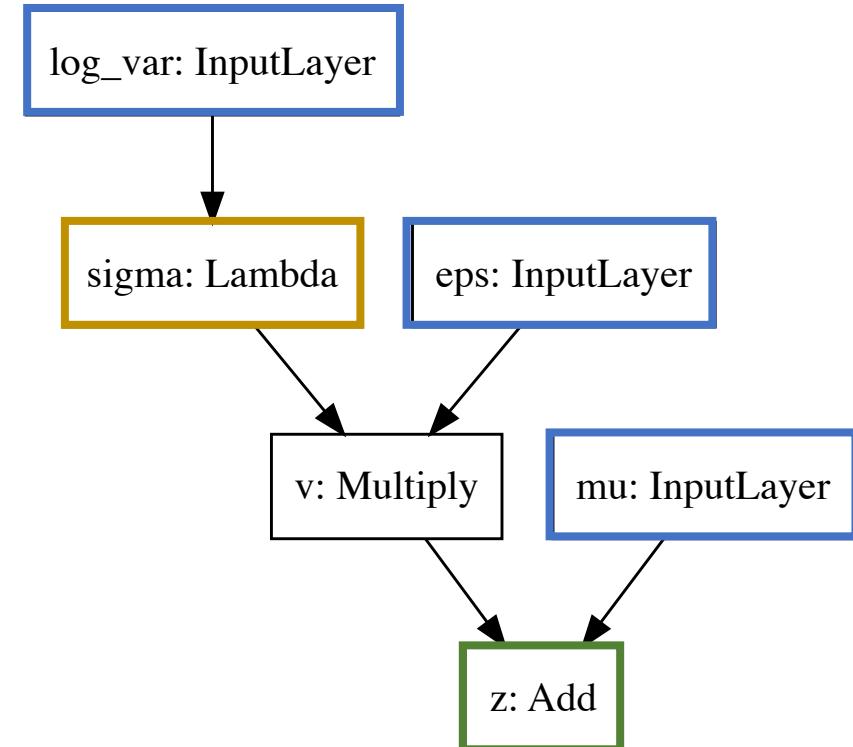


## 2. The Sampling Network



## 2. The Sampling Network

- Latent dimension:  $l$  ( $= 2$  in our implementation).
- Input of the sampling network.
  - Log variance:  $\eta$  ( $l$ -dim vector).
  - Mean:  $\mu$  ( $l$ -dim vector).
  - $\epsilon$ : randomly sampled from  $\mathcal{N}(\mathbf{0}, \mathbf{I}_l)$ .
- The sampling network:
  - $\sigma = e^{0.5\eta}$  ( $l$ -dim vector).
  - $v = \sigma \circ \epsilon$  ( $l$ -dim vector).
  - $z = \mu + v$  ( $l$ -dim vector).
  - Output  $z$ .



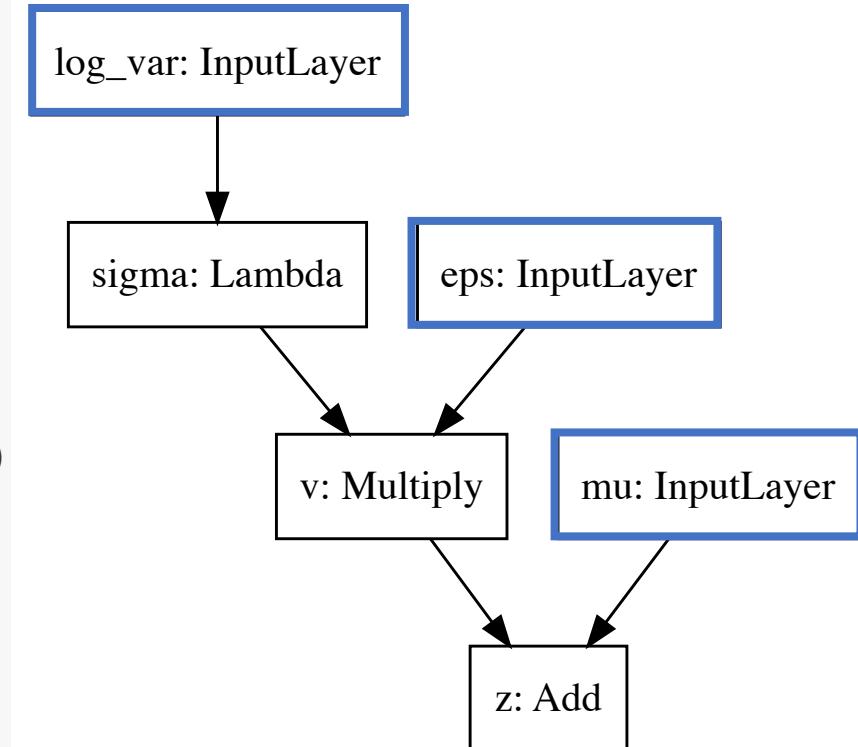
## 2. The Sampling Network

```
from keras.layers import Lambda, Multiply, Add
from keras import backend as K

# inputs
mu = Input(shape=(latent_dim,), name='mu')
log_var = Input(shape=(latent_dim,), name='log_var')
eps = Input(shape=(latent_dim,), name='eps')

# layers
sigma = Lambda(lambda t: K.exp(.5*t), name='sigma')(log_var)
v = Multiply(name='v')([sigma, eps])
z = Add(name='z')([mu, v])

# model
sampling = models.Model(inputs=[mu, log_var, eps],
                        outputs=z,
                        name='sampling')
```



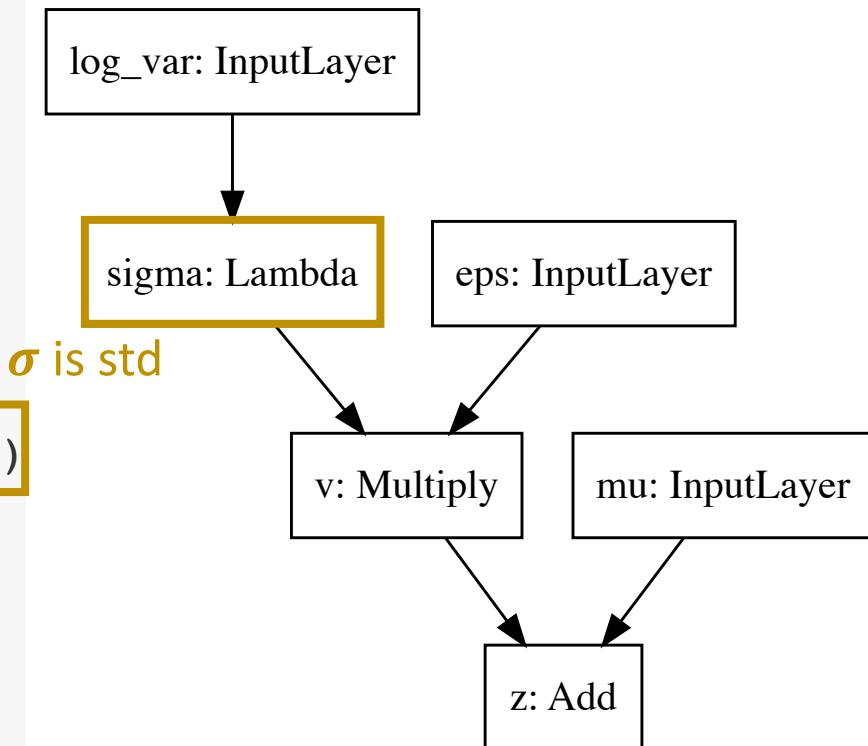
## 2. The Sampling Network

```
from keras.layers import Lambda, Multiply, Add
from keras import backend as K

# inputs
mu = Input(shape=(latent_dim,), name='mu')
log_var = Input(shape=(latent_dim,), name='log_var')
eps = Input(shape=(latent_dim,), name='eps')

# layers
sigma = Lambda(lambda t: K.exp(.5*t), name='sigma')(log_var)
v = Multiply(name='v')([sigma, eps])
z = Add(name='z')([mu, v])

# model
sampling = models.Model(inputs=[mu, log_var, eps],
                        outputs=z,
                        name='sampling')
```



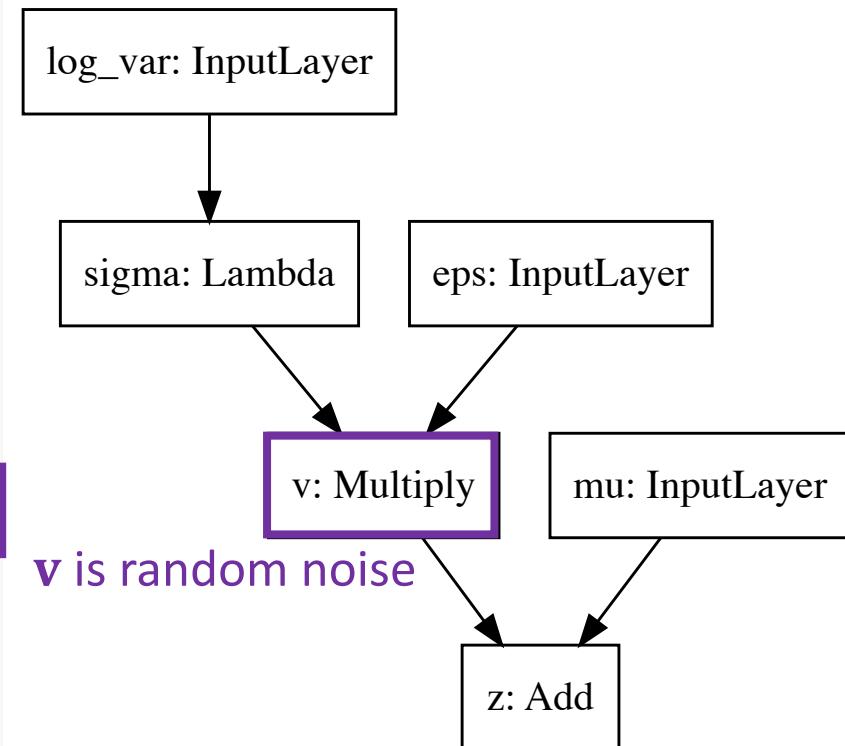
## 2. The Sampling Network

```
from keras.layers import Lambda, Multiply, Add
from keras import backend as K

# inputs
mu = Input(shape=(latent_dim,), name='mu')
log_var = Input(shape=(latent_dim,), name='log_var')
eps = Input(shape=(latent_dim,), name='eps')

# layers
sigma = Lambda(lambda t: K.exp(.5*t), name='sigma')(log_var)
v = Multiply(name='v')([sigma, eps])
z = Add(name='z')([mu, v])

# model
sampling = models.Model(inputs=[mu, log_var, eps],
                        outputs=z,
                        name='sampling')
```



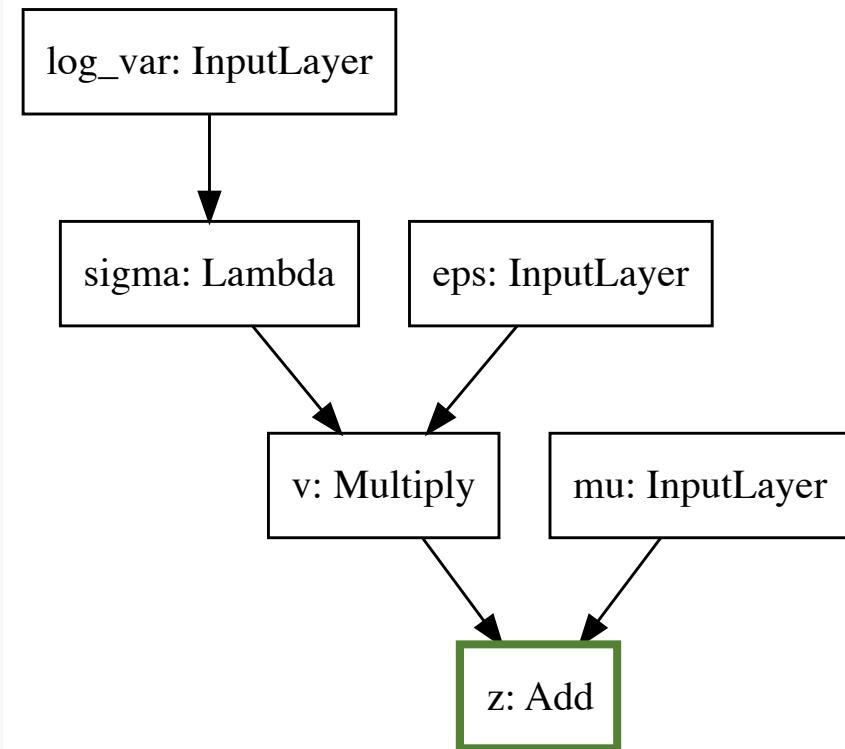
## 2. The Sampling Network

```
from keras.layers import Lambda, Multiply, Add
from keras import backend as K

# inputs
mu = Input(shape=(latent_dim,), name='mu')
log_var = Input(shape=(latent_dim,), name='log_var')
eps = Input(shape=(latent_dim,), name='eps')

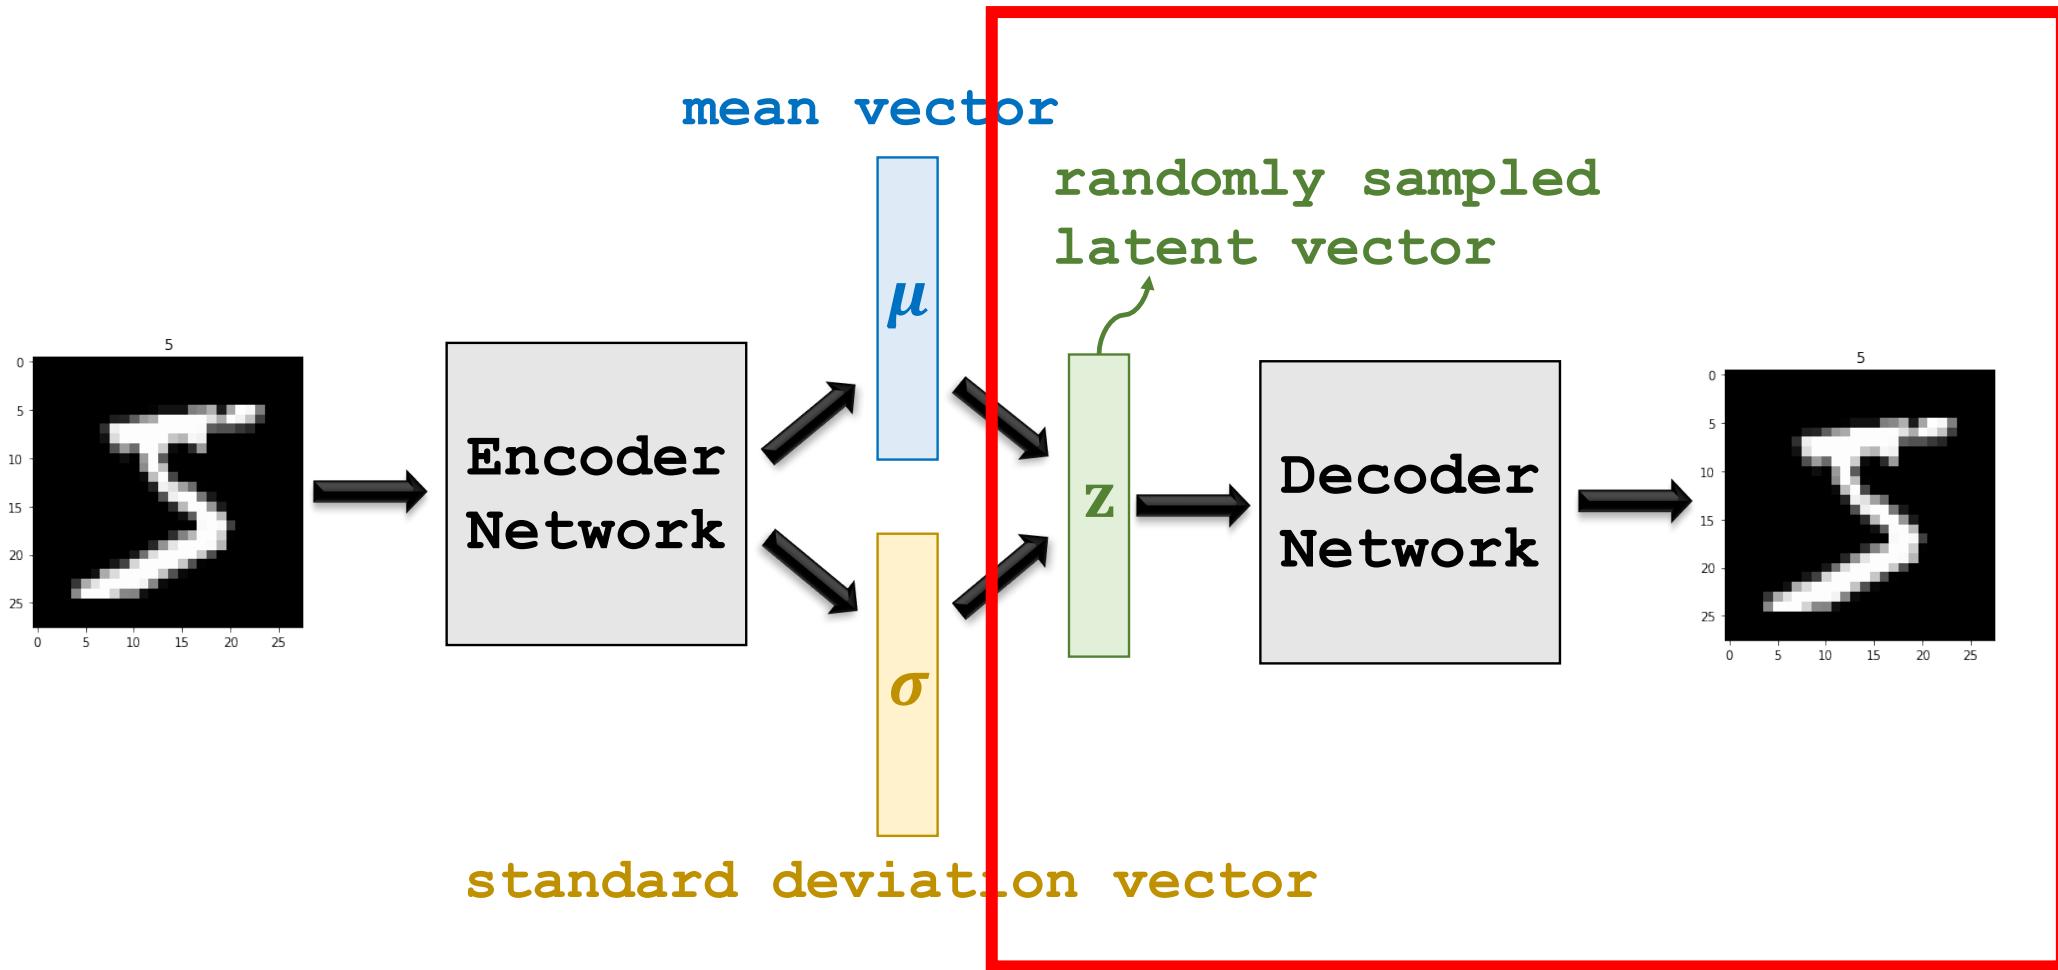
# layers
sigma = Lambda(lambda t: K.exp(.5*t), name='sigma')(log_var)
v = Multiply(name='v')([sigma, eps])
z = Add(name='z')([mu, v])

# model
sampling = models.Model(inputs=[mu, log_var, eps],
                        outputs=z,
                        name='sampling')
```



`z` is the sampled latent vector

### 3. The Decoder Network



### 3. The Decoder Network

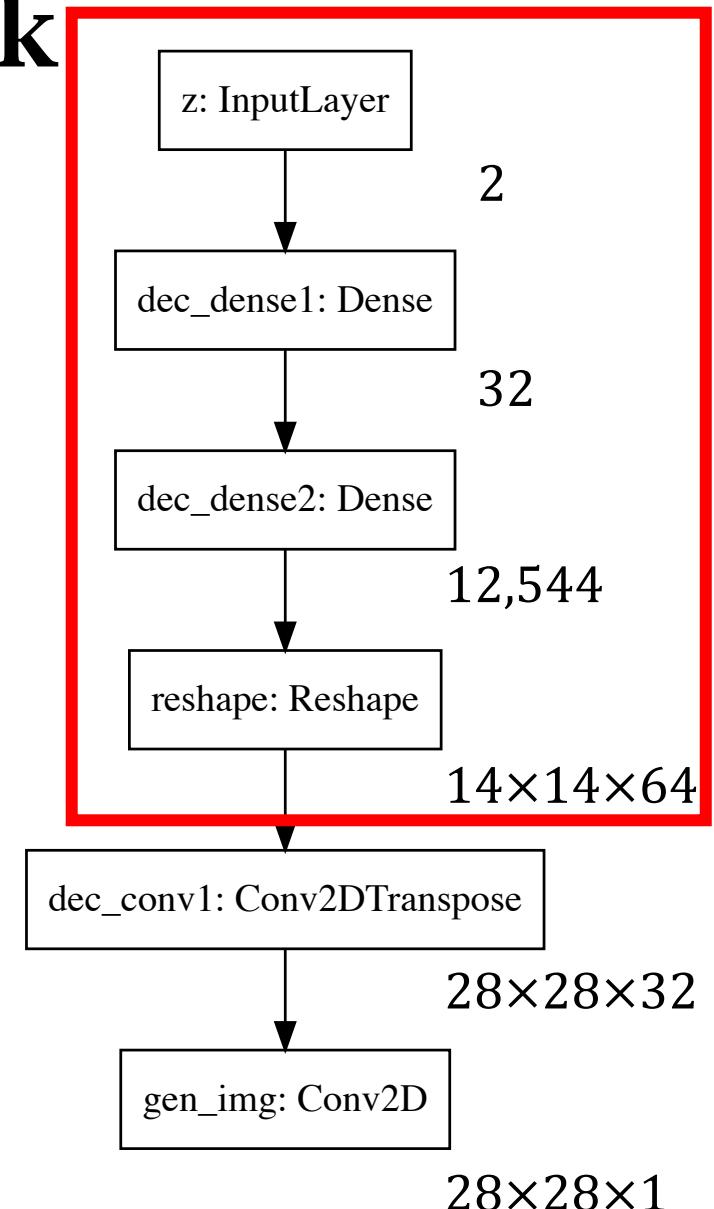
```
from keras import backend as K
import numpy

shape_before_flattening = K.int_shape(enc_conv_out4)[1:]
shape_after_flattening = numpy.prod(shape_before_flattening)

from keras.layers import Dense, Reshape, Conv2D, Conv2DTranspose

dec_dense1 = Dense(32, activation='relu', name='dec_dense1')
dec_dense2 = Dense(shape_after_flattening,
                   activation='relu', name='dec_dense2')
dec_reshape = Reshape(shape_before_flattening, name='reshape')

z = Input(shape=(latent_dim,), name='z')
dec_dense_out1 = dec_dense1(z)
dec_dense_out2 = dec_dense2(dec_dense_out1)
dec_reshape_out = dec_reshape(dec_dense_out2)
```

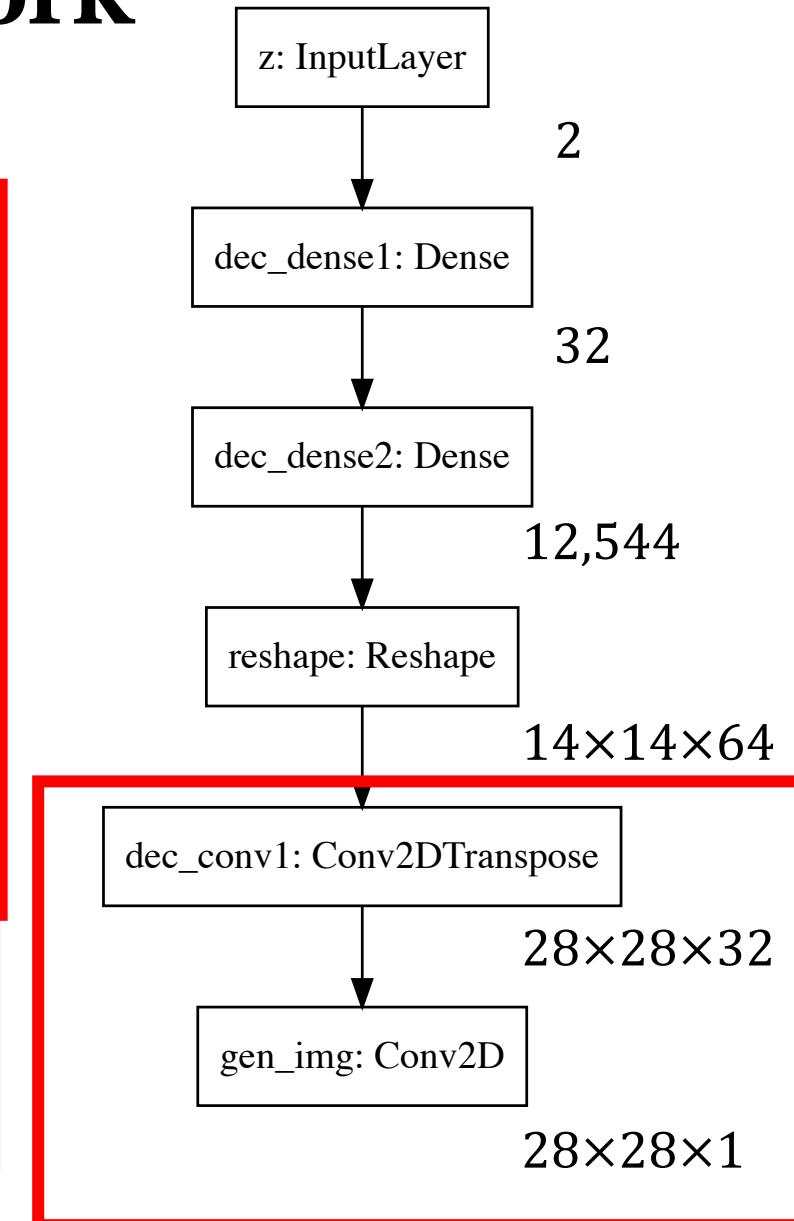


### 3. The Decoder Network

```
dec_conv1 = Conv2DTranspose(32, 3, padding='same',
                           activation='relu',
                           strides=(2, 2),
                           name='dec_conv1')
dec_conv2 = Conv2D(1, 3, padding='same',
                           activation='relu',
                           name='gen_img')

dec_conv_out1 = dec_conv1(dec_reshape_out)
gen_img = dec_conv2(dec_conv_out1)

decoder = models.Model(inputs=z,
                      outputs=gen_img,
                      name='decoder')
```

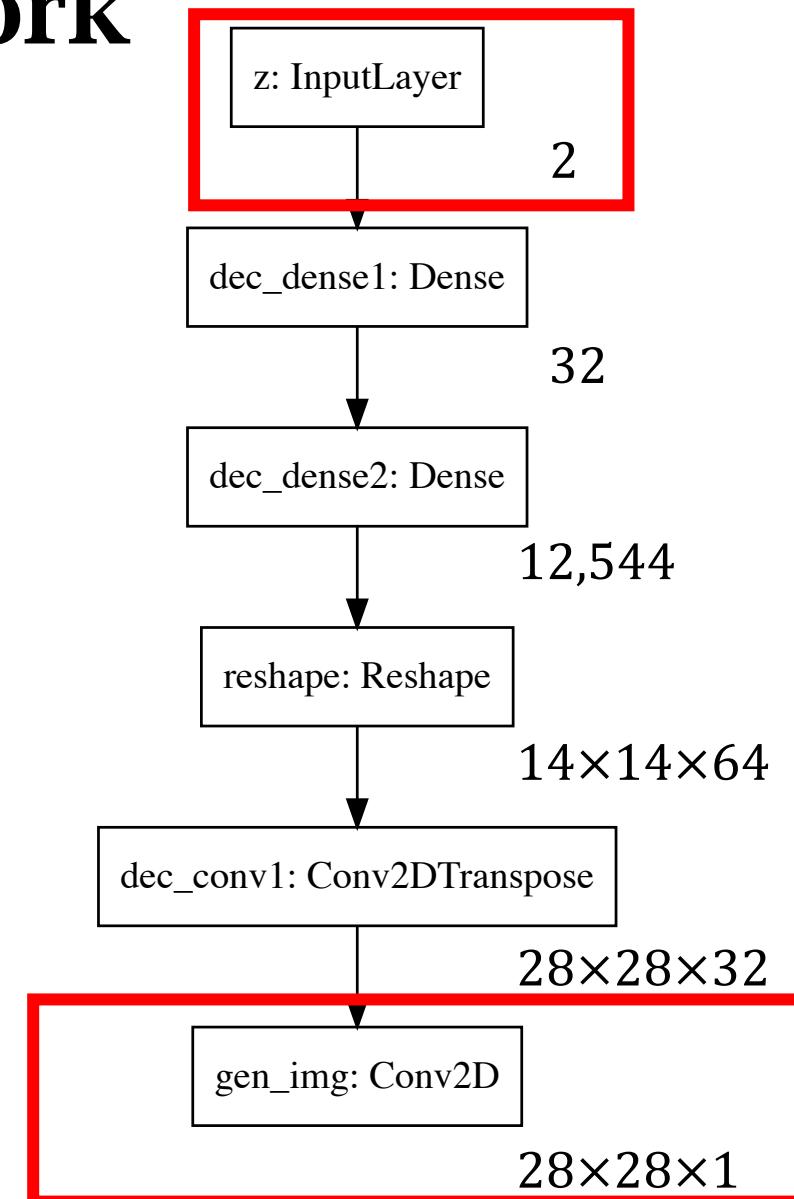


### 3. The Decoder Network

```
dec_conv1 = Conv2DTranspose(32, 3, padding='same',
                           activation='relu',
                           strides=(2, 2),
                           name='dec_conv1')
dec_conv2 = Conv2D(1, 3, padding='same',
                           activation='relu',
                           name='gen_img')

dec_conv_out1 = dec_conv1(dec_reshape_out)
gen_img = dec_conv2(dec_conv_out1)

decoder = models.Model(inputs=z,
                      outputs=gen_img,
                      name='decoder')
```



# 4. The Loss Function

- Generation loss:

$$\text{GenLoss} = \text{dist}(\text{input\_img}, \text{output\_img}).$$

```
l = keras.metrics.binary_crossentropy(input_img, output_img)
```

- Latent loss:

$$\text{LatLoss} = D_{\text{KL}}(p(\mathbf{z}) \parallel \mathcal{N}(0, \mathbf{I}))$$

```
l = -0.5 * K.mean(1 + log_var - K.square(mu) - K.exp(log_var), axis=-1)
```

- Loss = GenLoss +  $10^{-3} \cdot \text{LatLoss}$

The diagram illustrates the components of the total loss function. It shows the formula:  $\text{Loss} = \text{GenLoss} + 10^{-3} \cdot \text{LatLoss}$ . A purple arrow points from the term  $\text{GenLoss}$  to the text "Take input\_img and output\_img as input". Another purple arrow points from the term  $10^{-3} \cdot \text{LatLoss}$  to the text "tuning parameter". A third purple arrow points from the term  $\text{LatLoss}$  to the text "Take log\_var and mu as input".

Take `input_img` and `output_img` as input

tuning parameter

Take `log_var` and `mu` as input

# 4. The Loss Function

```
import keras
from keras import backend as K

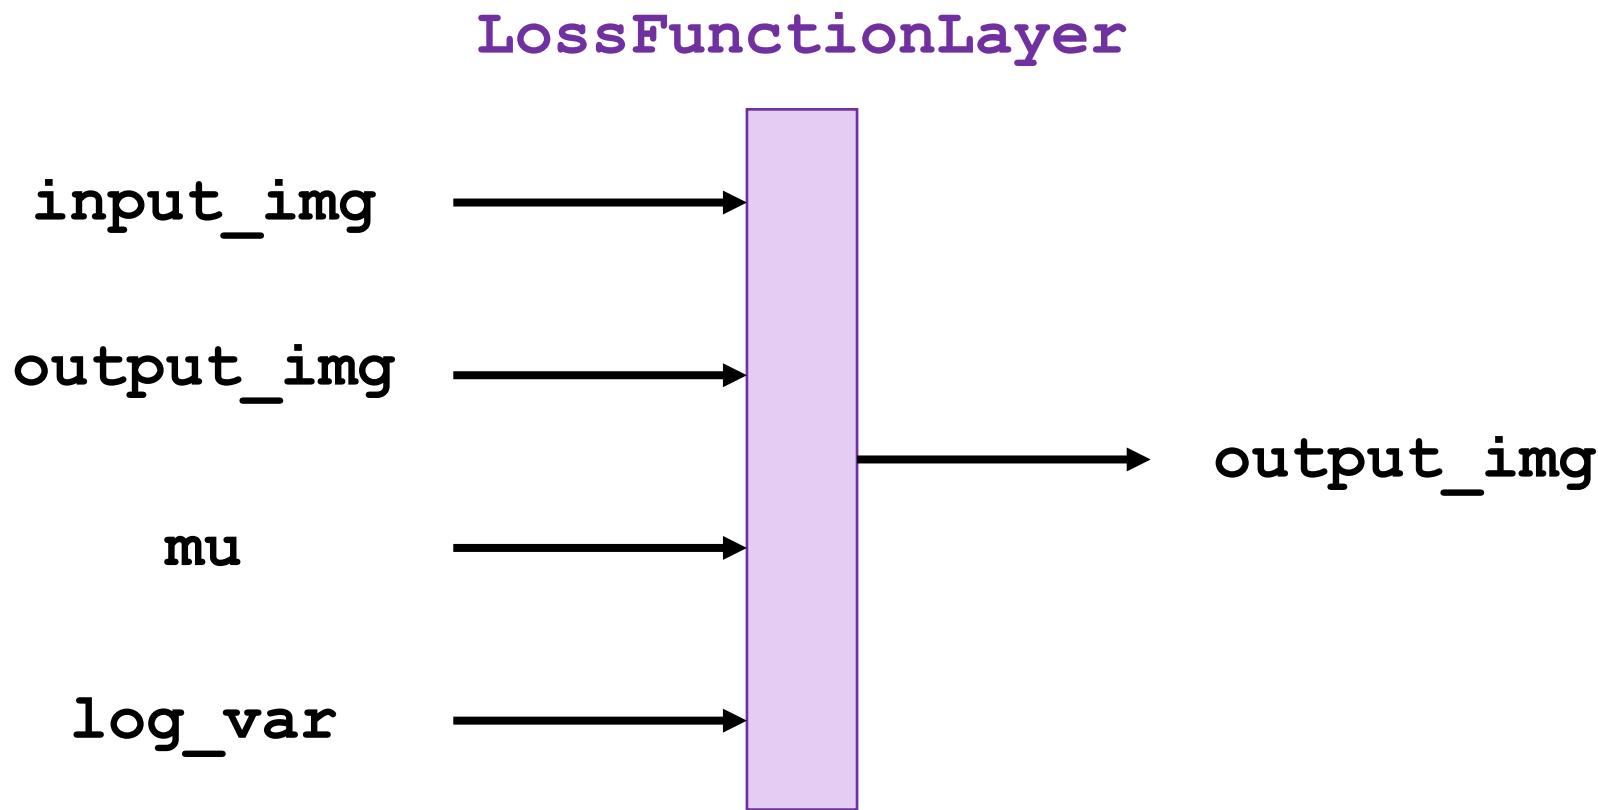
class LossFunctionLayer(keras.layers.Layer):
    lat_param = 1E-3

    def lat_loss(self, mu, log_var):
        l = -0.5 * K.mean(1 + log_var - K.square(mu) - K.exp(log_var), axis=-1)
        return self.lat_param * K.mean(l)

    def gen_loss(self, input_img, output_img):
        l = keras.metrics.binary_crossentropy(input_img, output_img)
        return K.mean(l)

    def call(self, inputs):
        input_img, output_img, mu, log_var = inputs
        gl = self.gen_loss(input_img, output_img)
        ll = self.lat_loss(mu, log_var)
        self.add_loss(gl+ll, inputs=inputs)
        return output_img
```

# 4. The Loss Function



Also compute the loss function.

- `input_img` and `output_img` for `GenerationLoss`.
- `mu` and `log_var` for `LatentLoss`.

# 5. Connect the Modules

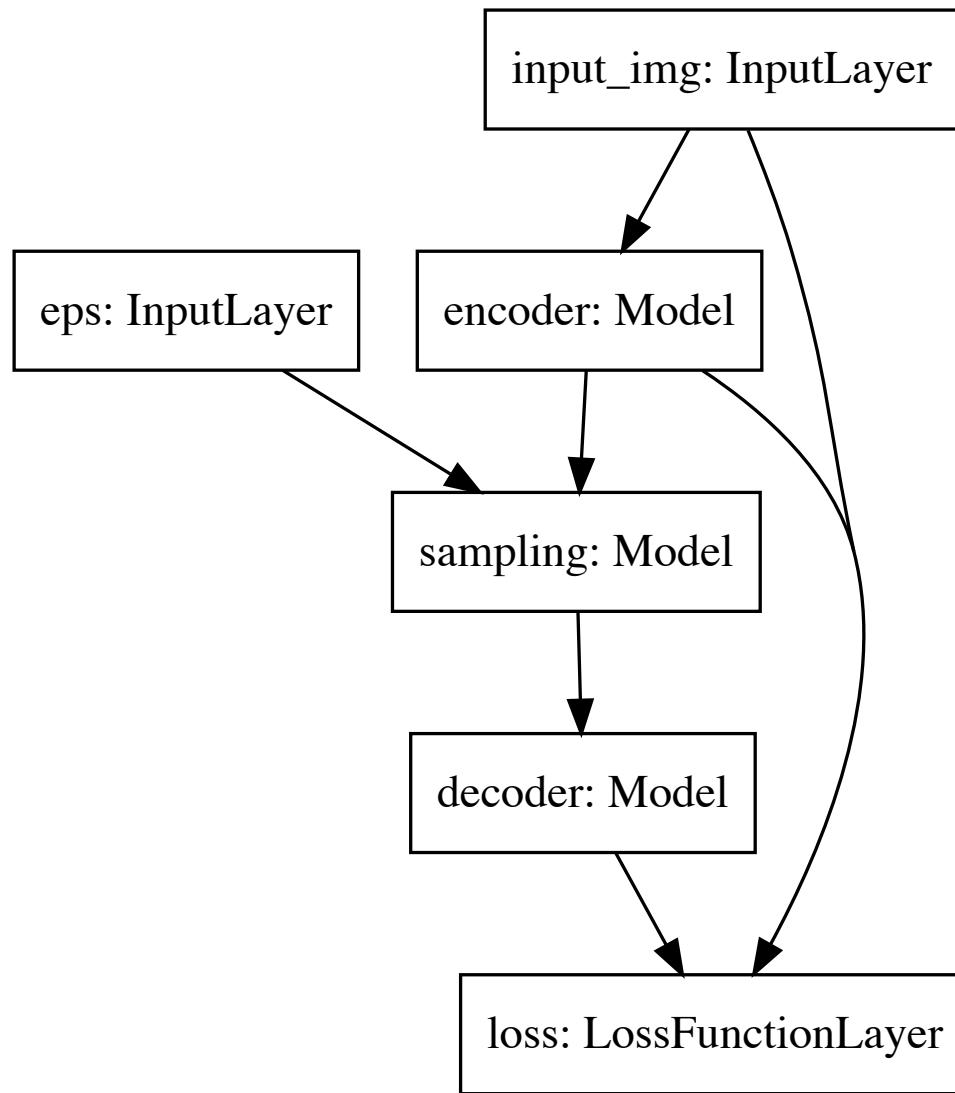
We have defined 4 modules:

- **Encoder**: [input\_img] → [mu, log\_var]
- **Sampling**: [mu, log\_var, eps] → [z]
- **Decoder**: [z] → [output\_img]
- **LossFunctionLayer**: [input\_img, output\_img, mu, log\_var] → [output\_img]

```
input_img = Input(shape=(28,28,1), name='input_img')
n = K.shape(input_img)[0]
eps = Input(tensor=K.random_normal(shape=(n, latent_dim)), name='eps')

mu, log_var = encoder(input_img)
z = sampling([mu, log_var, eps])
output_img = decoder(z)
output_img = LossFunctionLayer(name='loss')([input_img, output_img, mu, log_var])
model = models.Model(inputs=[input_img, eps], outputs=output_img)
```

# 5. Connect the Modules



# 6. Train the Model on the MNIST Dataset

```
history = model.fit(  
    x_train,  
    None,  
    shuffle=True,  
    epochs=50,  
    batch_size=128,  
    validation_data=(x_test, None)  
)
```

```
Train on 60000 samples, validate on 10000 samples  
Epoch 1/50  
60000/60000 [=====] - 231s 4ms/step - loss: 0.6717 - val_loss: 0.2194  
Epoch 2/50  
60000/60000 [=====] - 251s 4ms/step - loss: 0.2139 - val_loss: 0.2083  
Epoch 3/50  
60000/60000 [=====] - 299s 5ms/step - loss: 0.2066 - val_loss: 0.2096  
•  
•  
•  
Epoch 49/50  
60000/60000 [=====] - 240s 4ms/step - loss: 0.1839 - val_loss: 0.1859  
Epoch 50/50  
60000/60000 [=====] - 233s 4ms/step - loss: 0.1837 - val_loss: 0.1858
```

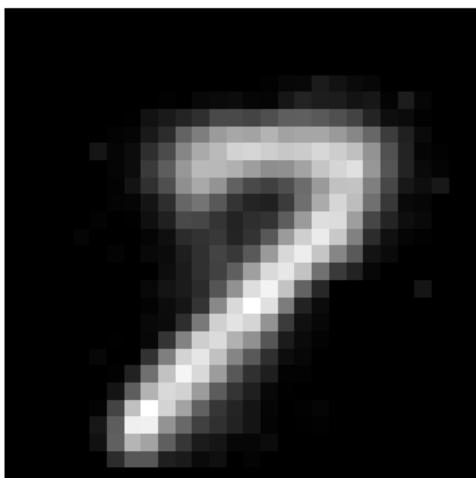
# 7. Visualize the Latent Space

arbitrary 2-dim vector

map the 2-dim vector to a  $28 \times 28$  image

```
z_sample = np.array([0.1, 0.2]).reshape((1, 2))
x_decoded = decoder.predict(z_sample)[0].reshape((28, 28))
```

```
fig = plt.figure(figsize=(6, 6))
plt.imshow(x_decoded, cmap='gray')
plt.axis('off')
plt.show()
```



# 7. Visualize the Latent Space

